

Applied Machine Learning

School of Computer Science, UPES, Dehradun.



LABORATORY FILE

B.TECH. - IV Semester

Jan - May. (2026)

SUBMITTED TO :-

Dr. Sahinur Rahman Laskar ,

Assistant Professor,

SoCS, UPES

SUBMITTED BY :-

Name : Pranav Joshi

Sap ID : 590012855

Batch : 19

Index

S.no	Experiment	Date
1	Data preprocessing	27-01-26
2	Predicting house price using linear regression	3-02-26
3	Predicting stock price using regression	10-02-26
4	Predicting customer churn rate using logistic regression	10-02-26
5	Spam detection (Classification and comparative analysis)	24-02-26
6	Credit Risk Assessment and comparative analysis	24-02-26
7	Anomaly detection and comparative analysis	24-02-26
8	Student Performance Level Analysis	17-02-26
9	Physiological signal classification	24-02-26
10	Iris classification	24-02-26
11	Prediction using bagging classifier	31-02-26
12	Regularization (ridge and lasso)	31-02-26
13	Mini Project	30-03-26

Experiment - 1

Aim => Text pre-processing in IPL dataset using python libraries .



Mount the drive

```
In [ ]: import pandas as pd
```

```
In [ ]: df=pd.read_csv("/content/drive/MyDrive/AIML_lab_sem4/expl/data.csv")
```

```
In [ ]: df.head()
```

```
Out[ ]:
```

	id	season	city	date	team1	team2	toss_winner	toss_d
0	1	2008	Bangalore	2008-04-18	Kolkata Knight Riders	Royal Challengers Bangalore	Royal Challengers Bangalore	
1	2	2008	Chandigarh	2008-04-19	Chennai Super Kings	Kings XI Punjab	Chennai Super Kings	
2	3	2008	Delhi	2008-04-19	Rajasthan Royals	Delhi Daredevils	Rajasthan Royals	
3	4	2008	Mumbai	2008-04-20	Mumbai Indians	Royal Challengers Bangalore	Mumbai Indians	
4	5	2008	Kolkata	2008-04-20	Deccan Chargers	Kolkata Knight Riders	Deccan Chargers	

```
In [ ]: import matplotlib.pyplot as plt
import seaborn as sns
```

```
In [ ]: df.columns
```

```
Out[ ]: Index(['id', 'season', 'city', 'date', 'team1', 'team2', 'toss_winner',
              'toss_decision', 'result', 'dl_applied', 'winner', 'win_by_runs',
              'win_by_wickets', 'player_of_match', 'venue', 'umpire1', 'umpire2',
              'umpire3'],
              dtype='object')
```

```
In [ ]: df.shape
```

```
Out[ ]: (577, 18)
```

Ques 1 . Team name normalization

1 Some team names appear in different textual formats (e.g., case differences or

extra spaces). Clean the team1 and team2 columns by:

2. Converting text to lowercase
3. Removing leading/trailing spaces
4. Verifying unique team names after cleaning

lower case all

```
In [ ]: df['team1'] = df['team1'].str.lower().str.strip()  
df['team2'] = df['team2'].str.lower().str.strip()
```

```
In [ ]: df[['team1', 'team2']]
```

```
Out[ ]:      team1      team2  
0  kolkata knight riders  royal challengers bangalore  
1  chennai super kings      kings xi punjab  
2    rajasthan royals      delhi daredevils  
3    mumbai indians  royal challengers bangalore  
4    deccan chargers      kolkata knight riders  
...      ...      ...  
572    delhi daredevils  royal challengers bangalore  
573    gujarat lions  royal challengers bangalore  
574  sunrisers hyderabad      kolkata knight riders  
575    gujarat lions      sunrisers hyderabad  
576  sunrisers hyderabad  royal challengers bangalore
```

577 rows x 2 columns

```
In [ ]: unique_teams = pd.concat([df['team1'], df['team2']]).unique()  
unique_teams
```

```
Out[ ]: array(['kolkata knight riders', 'chennai super kings', 'rajasthan royals',  
              'mumbai indians', 'deccan chargers', 'kings xi punjab',  
              'royal challengers bangalore', 'delhi daredevils',  
              'kochi tuskers kerala', 'pune warriors', 'sunrisers hyderabad',  
              'rising pune supergiants', 'gujarat lions'], dtype=object)
```

City Name Standardization

Standardize the city column by:

- 1) Replacing missing city names with "Unknown"
- 2) Converting all city names to title case
- 3) Counting matches played in each city

```
In [ ]: df.isnull().sum()
```

```
Out[ ]:
```

	0
id	0
season	0
city	7
date	0
team1	0
team2	0
toss_winner	0
toss_decision	0
result	0
dl_applied	0
winner	3
win_by_runs	0
win_by_wickets	0
player_of_match	3
venue	0
umpire1	0
umpire2	0
umpire3	577

dtype: int64

```
In [ ]: df['city'] = df['city'].fillna('Unknown')
```

```
In [ ]: df.isnull().sum()
```

```
Out[ ]:      0
         id    0
         season 0
         city   0
         date   0
         team1   0
         team2   0
         toss_winner 0
         toss_decision 0
         result   0
         dl_applied 0
         winner   3
         win_by_runs 0
         win_by_wickets 0
         player_of_match 3
         venue     0
         umpire1    0
         umpire2    0
         umpire3 577
```

dtype: int64

```
In [ ]: df['city'] = df['city'].str.strip().str.title()
```

```
In [ ]: df['city']
```

Out[]:

	city
0	Bangalore
1	Chandigarh
2	Delhi
3	Mumbai
4	Kolkata
...	...
572	Raipur
573	Bangalore
574	Delhi
575	Delhi
576	Bangalore

577 rows × 1 columns

dtype: object

```
In [ ]: city_counts = df['city'].value_counts()
city_counts
```

Out[]:

	count
city	
Mumbai	77
Bangalore	58
Kolkata	54
Delhi	53
Chennai	48
Chandigarh	42
Hyderabad	41
Jaipur	33
Pune	25
Durban	15
Centurion	12
Ahmedabad	12
Visakhapatnam	11
Dharamsala	9
Johannesburg	8
Unknown	7
Abu Dhabi	7
Cape Town	7
Port Elizabeth	7
Ranchi	7
Cuttack	7
Raipur	6
Sharjah	6
Rajkot	5
Kochi	5
Kimberley	3
East London	3
Nagpur	3
Bloemfontein	2
Indore	2

	count
city	
Kanpur	2

dtype: int64

3. Toss Decision Text Analysis

Analyze the toss_decision column:

- 1) Extract unique decisions
- 2) Count how many times each decision was taken
- 3) Visualize the frequency using a bar chart

```
In [ ]: choice = df['toss_decision'].unique()
choice
```

```
Out[ ]: array(['field', 'bat'], dtype=object)
```

```
In [ ]: dec = df['toss_decision'].value_counts()
dec
```


```
Out[ ]:
```

	count
toss_decision	
field	315
bat	262

dtype: int64

```
In [ ]: sns.barplot(x = choice , y = dec)
```

```
Out[ ]: <Axes: ylabel='count'>
```



4, Winner Name Extraction

From the winner column:

1) Identify and remove rows where the match result was "No Result" or "Tie"

2) Count how many matches each team won after cleaning text values

```
In [ ]: df_temp = df[df['result'] != 'no result']  
df_temp = df_temp[df_temp['result'] != 'tie']
```

```
In [ ]: winner_counts = df['winner'].str.lower().str.strip().value_counts()  
print(winner_counts)
```

```
winner
mumbai indians      80
chennai super kings  79
royal challengers bangalore  70
kolkata knight riders  68
rajasthan royals     63
kings xi punjab      63
delhi daredevils     56
sunrisers hyderabad  34
deccan chargers      29
pune warriors        12
gujarat lions        9
kochi tuskers kerala 6
rising pune supergiants 5
Name: count, dtype: int64
```

```
In [ ]: df.head()
```

```
Out[ ]:
   id  season  city  date  team1  team2  toss_winner  toss_de
0   1    2008  Bangalore  2008-04-18  kolkata knight riders  royal challengers bangalore  Royal Challengers Bangalore
1   2    2008  Chandigarh  2008-04-19  chennai super kings  kings xi punjab  Chennai Super Kings
2   3    2008    Delhi  2008-04-19  rajasthan royals  delhi daredevils  Rajasthan Royals
3   4    2008  Mumbai  2008-04-20  mumbai indians  royal challengers bangalore  Mumbai Indians
4   5    2008  Kolkata  2008-04-20  deccan chargers  kolkata knight riders  Deccan Chargers
```

5. Player of the Match Text Frequency

Perform text analysis on player of the match:

- 1) Remove null values
- 2) Find the top 10 most frequent player names
- 3) Plot the results using a Seaborn bar plot


```
In [ ]: pom_top_10 = df['player_of_match'].dropna().value_counts().head(10)
```

```
In [ ]: sns.barplot(x=pom_top_10.values, y=pom_top_10.index, palette='viridis')
plt.title('Top 10 Players of the Match')
plt.show()
```

/tmp/ipython-input-2314669101.py:1: FutureWarning:

Passing 'palette' without assigning 'hue' is deprecated and will be removed in v0.14.0. Assign the 'y' variable to 'hue' and set 'legend=False' for the same effect.

```
sns.barplot(x=pom_top_10.values, y=pom_top_10.index, palette='viridis')
```



6. Venue Tokenization

Count how many matches were played in each venue and plot a bar chart for the top 10.


```
In [ ]: venue_top_10 = df['venue'].value_counts().head(10)

sns.barplot(x=venue_top_10.values, y=venue_top_10.index, palette='magma')
plt.title('Top 10 Match Venues')
plt.show()
```

```
/tmp/ipython-input-3597930040.py:3: FutureWarning:
```

Passing 'palette' without assigning 'hue' is deprecated and will be removed in v0.14.0. Assign the 'y' variable to 'hue' and set 'legend=False' for the same effect.

```
sns.barplot(x=venue_top_10.values, y=venue_top_10.index, palette='magma')
```



7. Umpire Name Cleaning

Clean umpire columns (umpire1, umpire2, umpire3) by:

- 1) Replacing missing values with "Not Assigned"
- 2) Removing duplicate umpire names per match
- 3) Finding the most frequently officiating umpire

```
In [ ]: for col in ['umpire1', 'umpire2', 'umpire3']:
        df[col] = df[col].fillna('Not Assigned').str.strip()
```

```
In [ ]: all_umpires = pd.concat([df['umpire1'], df['umpire2'], df['umpire3']])
```

```
In [ ]: most_frequent = all_umpires[all_umpires != 'Not Assigned'].value_counts().idxmax
        print(f"Most frequently officiating umpire: {most_frequent}")
```

Most frequently officiating umpire: HDPK Dharmasena

8. Create a new text column match_summary by

combining:

team1, team2, winner, and season

Example: "MI vs CSK - MI won in 2019"

Display sample summaries.

```
In [ ]: def get_summary(row):  
        return f"{str(row['team1'])} vs {str(row['team2'])} at {str(row['venue'])}  
  
df['match_summary'] = df.apply(get_summary, axis=1)  
print(df['match_summary'].head())  
  
0    kolkata knight riders vs royal challengers ban...  
1    chennai super kings vs kings xi punjab at Punj...  
2    rajasthan royals vs delhi daredevils at Feroz ...  
3    mumbai indians vs royal challengers bangalore ...  
4    deccan chargers vs kolkata knight riders at Ed...  
Name: match_summary, dtype: object
```

9. Result Type Text Analysis

Analyze the result column:

Identify different textual result types


Count their occurrences

```
In [ ]: result_types = df['result'].value_counts()  
result_types
```

```
Out[ ]:      count  
result  
normal      568  
tie          6  
no result    3
```

dtype: int64

```
In [ ]: sns.countplot(data=df, x='result')
plt.title('Distribution of Match Result Types')
plt.show()
```



10. Toss Winner vs Match Winner (Text Matching)

Compare text values in `toss_winner` and `winner`:

Create a boolean column indicating whether the toss winner also won the match.
Visualize the comparison using a bar chart.

```
In [ ]: df['toss_winner'] = df['toss_winner'].str.lower().str.strip()
df['winner'] = df['winner'].str.lower().str.strip() # Normalize winner column
df['toss_match_winner'] = df['toss_winner'] == df['winner']

comparison = df['toss_match_winner'].value_counts()
comparison
```




```
Out[ ]:
```


	count
toss_match_winner	
True	291
False	286

dtype: int64

```
In [ ]: sns.barplot(x=comparison.index.map({True: 'Yes', False: 'No'}), y=comparison.v  
plt.title('Did the Toss Winner also win the Match?')  
plt.show()
```




```
In [ ]: winner_counts = df['winner'].str.lower().str.strip().value_counts()  
plt.figure(figsize=(12, 8))  
sns.barplot(x=winner_counts.values, y=winner_counts.index)  
  
plt.title('Winner Points Graph')  
plt.xlabel('No. of Wins')  
plt.ylabel('Team Name')  
plt.show()
```

```
In [ ]: import seaborn as sns
sns.barplot(x= df['winner'].value_counts().index, y=df['winner'].value_counts(
plt.xticks(rotation=45,ha='right')
```

```
Out[ ]: ([0, 1, 2, 3, 4, 5, 6, 7, 8, 9, 10, 11, 12],
[Text(0, 0, 'mumbai indians'),
Text(1, 0, 'chennai super kings'),
Text(2, 0, 'royal challengers bangalore'),
Text(3, 0, 'kolkata knight riders'),
Text(4, 0, 'rajasthan royals'),
Text(5, 0, 'kings xi punjab'),
Text(6, 0, 'delhi daredevils'),
Text(7, 0, 'sunrisers hyderabad'),
Text(8, 0, 'deccan chargers'),
Text(9, 0, 'pune warriors'),
Text(10, 0, 'gujarat lions'),
Text(11, 0, 'kochi tuskers kerala'),
Text(12, 0, 'rising pune supergiants')])
```



Experiment => 2

Predicting Housing Prices

Aim => Develop a regression model to predict house prices based on features like location, size, and amenities.

Problem Statement:

The determination of real estate market value is a complex process influenced by a diverse interplay of structural, locational, and economic variables . Historically, property valuation has relied heavily on manual appraisals and the subjective expertise of agents, which can often be inconsistent, time-consuming, or influenced by human bias. As the real estate market grows in complexity, there is an increasing demand for objective, data-driven solutions that can provide rapid and accurate valuations without the variance found in manual estimations.

This project addresses these inefficiencies by seeking to automate the valuation process through the identification of intricate patterns within historical housing data. Given a dataset containing 120 detailed property records, the primary objective is to develop a robust regression model capable of transforming property attributes into precise market estimates . By leveraging machine learning, we aim to move away from subjective "best guess" valuations toward a reliable, evidence-based framework that ensures transparency and mathematical consistency for both buyers and sellers .

- **Input Features :**

- **Location:** The neighborhood type (Rural, Urban, Suburban).
- **Area_sqft:** The total floor area of the house in square feet.
- **Bedrooms/Bathrooms:** The count of essential living spaces.
- **Balcony/Parking:** Additional amenities available.
- **Age_Years:** The age of the property (to account for depreciation or vintage value).
- **Furnished:** Status of the interior furnishing (Yes/No).

- **Target Variable :**

- **Price:** The estimated market value of the house (in thousands or millions of units).

Model Pipeline :

1. Data Understanding

- **Target Variable (y):** The target is to calculate the price of the house, which represents the numerical value we want the model to learn to predict .
- **Input Features :**
 - **Numerical:** Area_sqft, Bedrooms, Bathrooms, Age_Years, Parking, and Balconies.
 - **Categorical:** Location (Urban, Suburban, Rural) and Furnished (Yes, No).
- In the code, we will use Pandas functions like `pd.read_csv()` to load the dataset and `df.info()` to verify that each row represents an individual house record .

2. Data Preprocessing

- **Handling Categorical Variables:** Since machine learning models require numerical input, we will convert text based categories like "Urban" or "Suburban" into numbers . We will use One-Hot Encoding with help of `pandas.get_dummies()` to create binary columns like `Location_Urban` and `Furnished_Yes`, which avoids introducing false ordinal relationships .
- **Feature and target Split:** We separate the data into X which will contain all the input columns. and y which will have the output column which will be the price column.

3. Train–Test Split:

We use the `train_test_split()` function to divide the data into a Training set (80%) to learn the relationships and a Test set (20%) to evaluate performance .

4. Regression Model & Training

- **Model Selection:** We choose Linear Regression as the mathematical model to learn how each feature contributes to the final price .
- **Training Process:** Using the `model.fit()` function, the model calculates the Coefficients (weights) that minimize prediction error . This is achieved by minimizing the error on the training data to find the best fit line.

5. Predictions

- Once trained, the model is applied to the unseen test data to generate price estimates.
- We use the `model.predict()` function on the test features `X_test` to see how well the model generalizes to new information.
- We plot the actual prices against the predicted prices to visually confirm if the model is correctly following the market's upward or downward trends.

6. Evaluate the Model

1. We assess the model's accuracy using Mean Squared Error (MSE) and R-Square via the `sklearn.metrics` library .
2. A lower MSE indicates a better fit as it penalizes large errors strongly, while the R2 score tells us how much of the price variance is explained by our input features .

7. Saving the model

- To use this model in future applications without retraining, we save it as a file.
- We use the Pickle library (`pickle.dump()`) to serialize the trained model and save it as a `.pkl` file. This allows us to unpickle or load the model later using `pickle.load()` for instant predictions.



Predicting Housing Prices: Develop a regression model to predict house prices based on features like location, size, and amenities.

```
In [ ]: import pandas as pd
import matplotlib.pyplot as plt
```

```
In [ ]: df = pd.read_csv("/content/drive/MyDrive/AI/ML_lab_sem4/exp2,3,4/Housing_Price_
```

```
In [ ]: df.head()
```

```
Out[ ]:   HouseID  Location  Area_sqft  Bedrooms  Bathrooms  Balcony  Parking  Age_
0      H001    Rural    1028         1         3         0         0
1      H002    Rural    2316         1         4         0         0
2      H003    Urban    1834         1         2         2         2
3      H004    Urban    1719         3         1         0         2
4      H005  Suburban    1118         2         3         0         0
```

```
In [ ]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 120 entries, 0 to 119
Data columns (total 10 columns):
#   Column      Non-Null Count  Dtype
---  -
0   HouseID     120 non-null   object
1   Location    120 non-null   object
2   Area_sqft   120 non-null   int64
3   Bedrooms    120 non-null   int64
4   Bathrooms   120 non-null   int64
5   Balcony     120 non-null   int64
6   Parking     120 non-null   int64
7   Age_Years   120 non-null   int64
8   Furnished   120 non-null   object
9   Price       120 non-null   float64
dtypes: float64(1), int64(6), object(3)
memory usage: 9.5+ KB
```

```
In [ ]: df.isnull().sum()
```

```
Out[ ]:
0
HouseID 0
Location 0
Area_sqft 0
Bedrooms 0
Bathrooms 0
Balcony 0
Parking 0
Age_Years 0
Furnished 0
Price 0
```

dtype: int64

```
In [ ]: df = df.drop("HouseID" , axis = 1)
```

```
In [ ]: df = pd.get_dummies(df, columns=['Location', 'Furnished'], drop_first=True)

X = df.drop('Price', axis=1)
y = df['Price']
```

```
In [ ]: df.head()
```

```
Out[ ]:
Area_sqft  Bedrooms  Bathrooms  Balcony  Parking  Age_Years  Price  Locat
0      1028         1         3         0         0         5   55.84
1      2316         1         4         0         0         3  103.88
2      1834         1         2         2         2        23   86.42
3      1719         3         1         0         2        14   96.77
4      1118         2         3         0         0        13   63.94
```


```
In [ ]: from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
```

```
In [ ]: X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, randc
```

```
In [ ]: scaler = StandardScaler()
X_train_scaled = scaler.fit_transform(X_train)
X_test_scaled = scaler.transform(X_test)
```

```
In [ ]: from sklearn.linear_model import LinearRegression
```

```
In [ ]: lr = LinearRegression()  
lr.fit(X_train_scaled, y_train)
```

```
Out[ ]: 
LinearRegression()
```

```
In [ ]: y_pred = lr.predict(X_test_scaled)
y_pred
```

```
Out[ ]: array([101.39, 66.9, 63.94, 60.96, 62.13, 58.67, 35.35, 66.31,  
102.29, 63.55, 74.89, 114.76, 131.89, 77.38, 63.54, 113.85,  
87.52, 55.84, 104.41, 61.74, 129.28, 95.59, 101.56, 115.42])
```

```
In [ ]: lr.predict(X_test_scaled[0].reshape(1, -1))
```

```
Out[ ]: array([101.39])
```

```
In [ ]: price_correlations = df.corr(numeric_only=True)['Price'].sort_values(ascending  
display(price_correlations))
```

	Price
Price	1.000000
Area_sqft	0.717382
Location_Urban	0.459478
Bathrooms	0.333676
Bedrooms	0.154453
Parking	0.066709
Furnished_Yes	-0.046588
Balcony	-0.083178
Location_Suburban	-0.143702
Age_Years	-0.440987

dtype: float64

```
In [ ]: lr.intercept_
```

```
Out[ ]: np.float64(85.26947916666667)
```

```
In [ ]: lr.coef_
```



```
Out[ ]: array([ 1.50871230e+01,  5.49398938e+00,  7.47031161e+00, -4.21884749e-15,
                3.98258743e+00, -8.63680352e+00,  4.78260012e+00,  1.13632418e+01,
                -2.83106871e-15])
```

```
In [ ]: from sklearn.metrics import mean_squared_error, r2_score
```

```
In [ ]: mse = mean_squared_error(y_test, y_pred)
mse
```

```
Out[ ]: 2.6716089323484935e-28
```

```
In [ ]: r2 = r2_score(y_test, y_pred)
r2
```

```
Out[ ]: 1.0
```

```
In [ ]: import matplotlib.pyplot as plt
plt.plot(y_test.values, label='Actual Price', color='blue', marker='o')
plt.plot(y_pred, label='Predicted Price', color='red', linestyle='--', markers=
plt.title('Actual vs Predicted House Prices')
plt.xlabel('Test Case Number')
plt.ylabel('Price')
plt.legend()
plt.grid(True)
plt.savefig('actual_vs_predicted.png')
```



```
In [ ]: import pickle
```

```
In [ ]: folder_path = '/content/drive/MyDrive/AIML_lab_sem4/exp2,3,4'

# 1. Save the Model
with open(folder_path + 'house_model.pkl', 'wb') as f:
    pickle.dump(lr, f)

# 2. Save the Scaler (Don't forget this!)
with open(folder_path + 'scaler.pkl', 'wb') as f:
    pickle.dump(scaler, f)

print("Model and Scaler saved successfully to Google Drive!")
```

Model and Scaler saved successfully to Google Drive!

```
In [ ]: import pickle
import pandas as pd

folder_path = '/content/drive/MyDrive/AIML_lab_sem4/exp2,3,4'

# Load the saved model
with open(folder_path + 'house_model.pkl', 'rb') as f:
    loaded_model = pickle.load(f)

# Load the saved scaler
with open(folder_path + 'scaler.pkl', 'rb') as f:
    loaded_scaler = pickle.load(f)

print("Model and Scaler loaded successfully!")
```

Model and Scaler loaded successfully!

Now, let's create some example 'new' data. It's important that this new data has the same features and column order as your original training data (`X_train`).

```
In [ ]: new_house_data = pd.DataFrame({
    'Area_sqft': [2000],
    'Bedrooms': [3],
    'Bathrooms': [2],
    'Balcony': [1],
    'Parking': [1],
    'Age_Years': [10],
    'Location_Suburban': [False],
    'Location_Urban': [True],
    'Furnished_Yes': [True]
})

display(new_house_data)
```

	Area_sqft	Bedrooms	Bathrooms	Balcony	Parking	Age_Years	Location_Subu
0	2000	3	2	1	1	10	

Next, we need to scale this new data using the *loaded scaler* (not a new one). This is crucial for consistent predictions.

```
In [ ]: new_house_data_scaled = loaded_scaler.transform(new_house_data)
        predicted_price = loaded_model.predict(new_house_data_scaled)

        print(f"The predicted price for the new house is: ${predicted_price[0]:.2f}")
```

The predicted price for the new house is: \$112.00

This process ensures that your model always receives input data in the format it was trained on, leading to reliable predictions.

Key takeaways :

- The linear regression model successfully learned the relationship between housing features such as location, area, amenities, and age with the house price, producing stable and consistent predictions.
- The model achieved very low mean squared error and a high R^2 score, indicating strong predictive performance. This high accuracy is expected since the dataset is synthetic and well-structured.
- One-hot encoding of categorical variables and standardization of numerical features helped ensure that all input features contributed fairly to the regression model.
- Saving both the trained model and the scaler ensures the model can be reliably reused for future predictions on new housing data.

Experiment 3

Stock price prediction

AIM : Develop a time series prediction model to forecast stock prices.

Problem Overview :

The stock market is a dynamic environment where price determination is driven by a complex interplay of historical trends, market sentiment, and macroeconomic shifts. Traditional manual forecasting methods often struggle to maintain consistency, as human analysts can be biased or overwhelmed by the sheer volume of sequential data points . This project aims to address these challenges by developing an automated valuation system that identifies hidden temporal patterns within historical stock records.

By analyzing a dataset of stock records, the task is to build a robust regression model that can predict future market values based on past performance. Moving away from subjective estimates, this project implements a data-driven framework where the model learns from historical "lags" to forecast future outcomes. This ensures that the forecasting process is grounded in empirical evidence, providing a consistent and transparent tool for investors and analysts to anticipate market movements without the pitfalls of manual estimation .

Model Pipeline :

1. Understand the Dataset

- **Temporal Data:** Each row represents a specific point in time, tracking the movement of various stocks over a sequential period .
- **Target Variable (y):**
 - Price : The continuous value we want the model to forecast for the current day .
- **Input Features (X):**
 - **Numerical (Lags):** We use the prices from the previous 5 days (lag1 to lag5) as our primary predictors.
 - **Categorical:** The Stock ticker (e.g., AAPL) identifying which specific company we are forecasting.
- **Excluded Features:** We intentionally exclude same-day Open, High, Low, Close, and Volume to prevent the model from cheating and to ensure it relies only on historical trends.

2. Data Preprocessing

- **Chronological Alignment & Encoding:**
 - Machine learning models require numerical data and proper ordering, so we convert the Date column using `pd.to_datetime()` and sort the records by time.
 - We use One-Hot Encoding (via `pd.get_dummies()`) to convert stock names into binary numbers, allowing the model to distinguish between different stocks without assuming one is "higher" than another .
- **Lag Feature Engineering:**
 - To capture the "Time Series" behavior, we create 5 new columns using a `shift()` operation. These Lag Features allow the model to look at the prices from the last 5 days to predict today's value.
- **Feature-Target Split:**
 - We separate the data into X (the lags and stock dummies) and Y (the target Price) to define our forecasting function .

3. Train-Test Split

- Because this is time-series data, we do not shuffle the records. We must train on the "past" to predict the "future".
- We take the first 80% of the historical timeline for the Training set and reserve the final 20% for the Test set to evaluate true forecasting accuracy .

4. Choose and Train the Model

- We choose a Linear regression model to identify the mathematical trend line that connects past prices to future values .
- The model uses the `.fit()` function to calculate weights (coefficients) for each lag . It learns exactly how much weight to give to "yesterday's price" versus "the price 5 days ago" to minimize the prediction error.

5. Predictions

- The trained model is applied to the unseen test set using the `.predict()` function.
- We plot the actual prices against the predicted prices to visually confirm if the model is correctly following the market's upward or downward trends.

6. Evaluate the Model

- We evaluate the performance using Mean Absolute Error (MAE) and R-Squared (R2) .
- A lower MAE indicates the predicted prices are very close to reality, while the R2 score shows how much of the stock's volatility is explained by its own history .

MAE	1.27668
R2	1.00

7. Save the Model

- To use the forecaster in the future without retraining on old data, we save the completed model.
- We use the Pickle library (`pickle.dump()`) to serialize the trained Linear Regression model into a .pkl file. This allows for instant loading and real-time forecasting in future sessions.

```
In [ ]: import pandas as pd
import matplotlib.pyplot as plt
from sklearn.preprocessing import LabelEncoder
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
from sklearn.metrics import mean_squared_error, r2_score
```

```
In [ ]: df = pd.read_csv("/content/drive/MyDrive/AIIML_lab_sem4/exp2,3,4/stock_price_wi
```

```
In [ ]: df.head()
```

```
Out[ ]:
      Date  Stock  Open  High  Low  Close  Volume  Price
0  2020-01-01  AAPL  173.80  174.75  173.68  174.43  8204212  174.286667
1  2020-01-02  AAPL  176.01  177.59  174.07  177.55  2766891  176.403333
2  2020-01-03  AAPL  177.57  178.37  176.62  176.71  5721339  177.233333
3  2020-01-06  AAPL  176.01  177.58  171.33  171.73  9242680  173.546667
4  2020-01-07  AAPL  171.85  172.20  170.69  170.82  4416664  171.236667
```

```
In [ ]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 6525 entries, 0 to 6524
Data columns (total 8 columns):
#   Column  Non-Null Count  Dtype
---  ---
0   Date    6525 non-null    object
1   Stock   6525 non-null    object
2   Open    6525 non-null    float64
3   High    6525 non-null    float64
4   Low     6525 non-null    float64
5   Close   6525 non-null    float64
6   Volume  6525 non-null    int64
7   Price   6525 non-null    float64
dtypes: float64(5), int64(1), object(2)
memory usage: 407.9+ KB
```

```
In [ ]: df['Date'] = pd.to_datetime(df['Date'])
df = df.sort_values('Date')
df = df.reset_index(drop=True)
```

```
In [ ]: df_encoded = pd.get_dummies(df, columns=['Stock'], drop_first=True)
```

```
In [ ]: def create_lag_features(data, lag=5):
    df_lag = data.copy()
    stock_cols = [col for col in data.columns if col.startswith('Stock_')]

    for i in range(1, lag + 1):
        df_lag[f'lag_{i}'] = (df_lag.groupby(stock_cols)['Price'].shift(i))
```



```
    return df_lag

df_lag = create_lag_features(df_encoded, lag=5)
df_lag.dropna(inplace=True)
```

```
In [ ]: X = df_lag.drop(['Date', 'Price'], axis=1)
        y = df_lag['Price']

        split_index = int(len(df_lag) * 0.8)

        X_train, X_test = X[:split_index], X[split_index:]
        y_train, y_test = y[:split_index], y[split_index:]
```

```
In [ ]: from sklearn.linear_model import LinearRegression

        model = LinearRegression()
```

```
In [ ]: model.fit(X_train, y_train)
```

```
Out[ ]: + LinearRegression
        LinearRegression()
```

```
In [ ]: y_pred = model.predict(X_test)
        y_pred
```

```
Out[ ]: array([194.79333333, 226.06333333, 246.99      , ..., 216.85333333,
               222.22      , 218.07      ])
```

```
In [ ]: from sklearn.metrics import mean_absolute_error, r2_score

        print("MAE :", mean_absolute_error(y_test, y_pred))
        print("R2  :", r2_score(y_test, y_pred))
```

```
MAE : 1.2766876649878284e-10
R2  : 1.0
```

```
In [ ]: import matplotlib.pyplot as plt
        import numpy as np


        plt.figure(figsize=(6, 6))

        # Scatter: actual vs predicted
        plt.scatter(y_test, y_pred)

        # Ideal linear line (y = x)
        min_val = min(y_test.min(), y_pred.min())
        max_val = max(y_test.max(), y_pred.max())
        plt.plot([min_val, max_val], [min_val, max_val])

        plt.xlabel("Actual Price")
        plt.ylabel("Predicted Price")
```

```
plt.title("Actual vs Predicted (Linear Regression Behavior)")
plt.show()
```



```
In [ ]: corr = df.corr(numeric_only=True)
display(corr)
```

	Open	High	Low	Close	Volume	Price
Open	1.000000	0.999589	0.999579	0.999010	0.002314	0.999543
High	0.999589	1.000000	0.999485	0.999576	0.001795	0.999837
Low	0.999579	0.999485	1.000000	0.999590	0.002654	0.999842
Close	0.999010	0.999576	0.999590	1.000000	0.001975	0.999872
Volume	0.002314	0.001795	0.002654	0.001975	1.000000	0.002141
Price	0.999543	0.999837	0.999842	0.999872	0.002141	1.000000

```
In [ ]: import pickle
```

```
In [ ]: with open('/content/drive/MyDrive/AIML_lab_sem4/exp2,3,4/stocks_model.pkl', 'w') as file:  
        pickle.dump(model, file)  
  
        print("Model saved successfully to '/content/drive/MyDrive/AIML_lab_sem4/exp2,3,4/stocks_model.pkl'")
```

Experiment = 4

Customer churn prediction

AIM : Customer Churn Prediction: Develop a model to predict customer churn in a subscription-based business.

Problem Statement :

Customer retention is a major challenge for service-based businesses, as acquisition costs often exceed retention costs. Manual tracking typically fails to identify the subtle behavioral patterns that signal a "churn" event. This project aims to automate the identification of at-risk customers by analyzing a dataset of 200 records including historical usage and contract data.

Unlike prior regression tasks, this is a Classification problem focused on predicting a discrete category: Churn (1) or No Churn (0). Implementing this data-driven framework allows the business to transition from reactive to proactive retention strategies. By targeting interventions toward the users most likely to depart, the business can optimize marketing resources and stabilize long-term revenue.

Project Pipeline :

1. Data Understanding

The dataset consists of 200 customer records with 11 attributes.

- Target Variable (y): Churn.
- Numerical Features: Age, Tenure_Months, MonthlyCharges, TotalCharges, and SupportTickets.
- Categorical Features: Gender, SubscriptionType, PaymentMethod, and ContractType.

2. Data Preprocessing

Raw data is cleaned and transformed to make it compatible with the Logistic Regression algorithm.

- Feature Selection: The CustomerID column is dropped using `df.drop()` as it is a unique identifier with no predictive value.
- Cleaning: Rows with missing values are removed via `df.dropna()`.
- Categorical Encoding: Text-based features are converted into numerical binary vectors using One-Hot Encoding (`pd.get_dummies()`).

3. Feature-Target & Train-Test Split

The processed data is separated into features (X) and the target (y), then divided to ensure a fair evaluation.

- Ratio: 80% of the data is used for training, and 20% is reserved for testing.
- Test Sample: This results in 40 unseen records used to validate the model's accuracy.

4. Choose and Train the Model

Logistic Regression is selected for its simplicity and interpretability in binary tasks.

- Implementation: The model is initialized with `max_iter=1000` and trained using the `.fit()` method.

5. Make Predictions

The trained model applies its learned weights to the test set to flag potential churn risks.

- The `.predict()` function is used on `X_test`.
- An array of binary predictions (e.g., `array([0, 0, 1...])`), assigning a churn status to each customer in the test group.

6. Evaluate the Model

The performance is measured by comparing the model's guesses against the actual outcomes.

- Accuracy: 75% (The model correctly identified 30 out of 40 customers).
- Classification Report:
 - Class 0 (Stayed): Precision and Recall are high at approximately 85%.
 - Class 1 (Churned): Precision and Recall are low at 29%, indicating the model struggles to catch the minority churn class..

Accuracy	Precision	Recall	F1
0.75	0.85	0.85	0.85

7. Save the Model

The model is preserved so it can be deployed immediately in a business dashboard without retraining.

- We use the Pickle library (`pickle.dump()`) to serialize the model object.
- Output: A permanent file named `customerchurn` is saved in the Drive for instant future use.

```
In [1]: import pandas as pd
        from sklearn.model_selection import train_test_split
        from sklearn.linear_model import LogisticRegression
        from sklearn.ensemble import RandomForestClassifier
        from sklearn.metrics import accuracy_score, classification_report
        from sklearn.preprocessing import StandardScaler
```

```
In [2]: df = pd.read_csv('/content/drive/MyDrive/AIML_lab_sem4/exp2,3,4/customer_churn.csv')
```

```
In [3]: df.head()
```

```
Out[3]:
```

	CustomerID	Gender	Age	Tenure_Months	SubscriptionType	MonthlyCharges
0	CUST1000	Male	49	54	Standard	99.33
1	CUST1001	Female	56	8	Standard	12.01
2	CUST1002	Male	66	27	Basic	57.62
3	CUST1003	Male	69	27	Premium	97.08
4	CUST1004	Male	49	34	Standard	54.65

```
In [4]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 200 entries, 0 to 199
Data columns (total 11 columns):
#   Column              Non-Null Count  Dtype
---  -
0   CustomerID          200 non-null    object
1   Gender              200 non-null    object
2   Age                 200 non-null    int64
3   Tenure_Months       200 non-null    int64
4   SubscriptionType     200 non-null    object
5   MonthlyCharges      200 non-null    float64
6   TotalCharges        200 non-null    float64
7   PaymentMethod       200 non-null    object
8   SupportTickets      200 non-null    int64
9   ContractType        200 non-null    object
10  Churn               200 non-null    int64
dtypes: float64(2), int64(4), object(5)
memory usage: 17.3+ KB
```

```
In [5]: df = df.drop(columns=['CustomerID'])
```

```
In [6]: df = df.dropna()
```

```
In [7]: df = pd.get_dummies(df, drop_first=True)
```

```
In [8]: from sklearn.model_selection import train_test_split
```

```
In [9]: X = df.drop('Churn', axis=1)
        y = df['Churn']
```

```
In [10]: X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)
```

```
In [11]: from sklearn.linear_model import LogisticRegression

        model = LogisticRegression(max_iter=1000)
```

```
In [12]: model.fit(X_train, y_train)
```

/usr/local/lib/python3.12/dist-packages/sklearn/linear_model/_logistic.py:465:
ConvergenceWarning: lbfgs failed to converge (status=1):
STOP: TOTAL NO. OF ITERATIONS REACHED LIMIT.

Increase the number of iterations (max_iter) or scale the data as shown in:
<https://scikit-learn.org/stable/modules/preprocessing.html>
Please also refer to the documentation for alternative solver options:
https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression

```
Out[12]: LogisticRegression(max_iter=1000)
```

```
In [13]: y_pred = model.predict(X_test)
```

```
In [14]: y_pred
```

```
Out[14]: array([0, 0, 0, 0, 1, 0, 0, 0, 1, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 1, 0, 0, 1, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0])
```

```
In [15]: from sklearn.metrics import accuracy_score, classification_report, confusion_matrix

        print("Accuracy:", accuracy_score(y_test, y_pred))
        print(classification_report(y_test, y_pred))
```

Accuracy: 0.75


	precision	recall	f1-score	support
0	0.85	0.85	0.85	33
1	0.29	0.29	0.29	7
accuracy			0.75	40
macro avg	0.57	0.57	0.57	40
weighted avg	0.75	0.75	0.75	40

```
In [17]: import seaborn as sns
        from sklearn.metrics import confusion_matrix
        import numpy as np
```



```
import matplotlib.pyplot as plt

cm = confusion_matrix(y_test, y_pred)
plt.figure(figsize=(8, 6))
sns.heatmap(cm, annot=True, fmt='d', cmap='Blues', cbar=False,
            xticklabels=['Predicted 0', 'Predicted 1'],
            yticklabels=['Actual 0', 'Actual 1'])
plt.xlabel('Predicted Label')
plt.ylabel('True Label')
plt.title('Confusion Matrix')
plt.show()
```



```
In [18]: true_correct = np.sum((y_test == 1) & (y_pred == 1))
true_wrong  = np.sum((y_test == 1) & (y_pred == 0))
false_correct = np.sum((y_test == 0) & (y_pred == 0))
false_wrong  = np.sum((y_test == 0) & (y_pred == 1))


labels = ['Churn = Yes (1)', 'Churn = No (0)']
correct = [true_correct, false_correct]
wrong = [true_wrong, false_wrong]

x = np.arange(len(labels))
width = 0.35
```



```
plt.figure(figsize=(6, 4))
plt.bar(x - width/2, correct, width, label='Correct Predictions')
plt.bar(x + width/2, wrong, width, label='Wrong Predictions')

plt.xticks(x, labels)
plt.ylabel("Number of Customers")
plt.title("Correct vs Wrong Predictions by Churn Class")
plt.legend()
plt.show()
```



```
In [19]: import pickle

model_filename = 'logistic_regression_model.pkl'
with open('/content/drive/MyDrive/AIML_lab_sen4/exp2,3,4/customerchurn', 'wb')
    pickle.dump(model, file)

print(f"Model successfully saved to {model_filename}")
Model successfully saved to logistic_regression_model.pkl
```

Experiment 5

Spam detection

Aim = Spam Email Detection: Build a spam email filter using text classification algorithms and perform comparative analysis. (Naive Bayes, Logistic Regression, Support Vector Machine, Decision Tree, Random Forest)

Problem Statement :

Spam detection is essential for maintaining email security and user productivity, as malicious or unsolicited emails can lead to security breaches or cluttered inboxes. Manual filtering is impossible given the volume of communication, necessitating automated classification models. This project aims to build a robust spam filter by performing a comparative analysis of five text classification algorithms: Naïve Bayes, Logistic Regression, Support Vector Machine (SVM), Decision Tree, and Random Forest. By evaluating these models, the goal is to identify which algorithm most accurately distinguishes between "Spam" and "Ham" (legitimate) emails, enabling proactive security and improved resource management.

Pipeline :

1. Problem Definition

The primary goal of this experiment is to develop a robust automated spam filter using advanced text classification techniques.

- **Classification Task:** The model must determine if an incoming email is Spam (1) or Ham (0).
- **Objective:** Perform a detailed comparative analysis across five distinct algorithms—Naïve Bayes, Logistic Regression, Support Vector Machine (SVM), Decision Tree, and Random Forest—to identify the most reliable architecture for real-world deployment.

2. Data Understanding

- **Dataset Content:** A collection of thousands of email text bodies labeled by human moderators as either spam or legitimate (ham).
- **Target (y):** Binary labels where 1 represents malicious or unsolicited "Spam" and 0 represents legitimate "Ham."
- **Input Features (X):** The raw text content of the emails, including the subject lines and message bodies.

3. Data Preprocessing

Since machine learning models cannot interpret raw text, the data is put through a rigorous cleaning and transformation process.

- Text Cleaning: Stripping out punctuation, special characters, and numbers to reduce noise and focus only on significant vocabulary.
- Stop-word Removal: Eliminating frequently used words (e.g., "and", "the", "is") that appear in almost every email and carry little to no predictive value.
- Vectorization: Utilizing TF-IDF (Term Frequency-Inverse Document Frequency) to transform words into numerical vectors. This method balances how often a word appears in one email against how common it is across the entire dataset.

4. Train-Test Split

To ensure the filter works on emails it has never seen before, the dataset is strategically divided.

- Data Allocation: The cleaned dataset is split into a Training set (80%) to build the models and a Test set (20%) to validate them.
- Purpose: This division prevents "overfitting" and allows us to measure how accurately the filter will perform in a live inbox environment.

5. Model Selection and training

Five different classification models are initialized and trained on the same training data to ensure a fair comparison:

1. Naïve Bayes: A probabilistic classifier based on Bayes' Theorem, highly efficient for high-dimensional text data.
2. Logistic Regression: A linear model that estimates the probability of an email belonging to the spam class based on the frequency of specific words.
3. Support Vector Machine (SVM): Finds the optimal hyperplane (boundary) that maximizes the distance between the spam and ham clusters.
4. Decision Tree: Uses a tree-like structure of word-based logical decisions to categorize each message.
5. Random Forest: An ensemble method that constructs a multitude of decision trees and outputs the class that is the mode of the individual trees, significantly improving stability.

6. Predictions

Once training is complete, each of the five models is tasked with classifying the 40 unseen emails in the test set.

- The `.predict()` function is called for each algorithm, generating five separate arrays of binary predictions.
- These predictions are compared directly against the actual ground-truth labels to pinpoint errors such as "False Positives" (blocking a real email) or "False Negatives" (letting spam through).

7. Model Comparison and evaluation :

The performance of each model is measured using Accuracy, Precision, and Recall. Based on the experimental outcomes.

	Accuracy	Precision	Recall
Naive Bayes	0.976	1.0	0.825
Logistic Regression	0.967	1.0	0.758
SVM	0.984	1.0	0.88
Decision Tree	0.968	0.9	0.859
Random Forest	0.979	1.0	0.83

Accuracy: Correct prediction rate of the model. SVM achieved the best accuracy (98.7%).

Precision: Measures how many predicted positives are actually correct. Logistic Regression, Naive Bayes, SVM, and Random Forest scored 1.

Recall: Measures how many actual positives are correctly identified. SVM achieved the highest recall (88%).

Overall SVM is the best model among the all as it scored hughes accuracy and recall as well as perfect precision .

8. Model saving

After identifying the superior architecture, the final step is ensuring the model is persistent.

- Implementation: We will save all the trained models using the Pickle library (`pickle.dump()`) to serialize the model objects into permanent files.
- Final Output: The generated `.pkl` files allow for immediate deployment in an email system for real-time filtering without the need for retraining.

```
In [2]: import pandas as pd
```

```
In [3]: df = pd.read_csv("/content/drive/MyDrive/AIML_lab_sem4/exp 5,6,7/spam_email.csv")
```

```
In [4]: df.head()
```

```
Out[4]:
```

	Category	Message
0	ham	Go until jurong point, crazy.. Available only ...
1	ham	Ok lar... Joking wif u oni...
2	spam	Free entry in 2 a wkly comp to win FA Cup fina...
3	ham	U dun say so early hor... U c already then say...
4	ham	Nah I don't think he goes to usf, he lives aro...

```
In [5]: import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn.feature_extraction.text import TfidfVectorizer
from sklearn.metrics import accuracy_score, classification_report, precision_s
```

```
In [6]: from sklearn.linear_model import LogisticRegression
from sklearn.svm import SVC
from sklearn.naive_bayes import MultinomialNB
from sklearn.tree import DecisionTreeClassifier
from sklearn.ensemble import RandomForestClassifier
```

```
In [7]: df['Category'] = df['Category'].map({'spam': 1, 'ham': 0})
print(df.head())
```

```
Category
```

	Category	Message
0	0	Go until jurong point, crazy.. Available only ...
1	0	Ok lar... Joking wif u oni...
2	1	Free entry in 2 a wkly comp to win FA Cup fina...
3	0	U dun say so early hor... U c already then say...
4	0	Nah I don't think he goes to usf, he lives aro...

```
In [8]: X = df['Message']
y = df['Category']
```

```
In [9]: X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, rando
```

```
In [10]: vectorizer = TfidfVectorizer(stop_words='english')
X_train_vec = vectorizer.fit_transform(X_train)
X_test_vec = vectorizer.transform(X_test)
```

```
In [11]: LR = LogisticRegression()
LR.fit(X_train_vec, y_train)
```

```
Out[11]: LogisticRegression
LogisticRegression()
```

```
In [12]: NB = MultinomialNB()
NB.fit(X_train_vec, y_train)
```

```
Out[12]: MultinomialNB
MultinomialNB()
```

```
In [13]: DT = DecisionTreeClassifier()
DT.fit(X_train_vec, y_train)
```

```
Out[13]: DecisionTreeClassifier
DecisionTreeClassifier()
```

```
In [14]: RF = RandomForestClassifier()
RF.fit(X_train_vec, y_train)
```

```
Out[14]: RandomForestClassifier
RandomForestClassifier()
```

```
In [15]: svc = SVC()
svc.fit(X_train_vec, y_train)
```

```
Out[15]: SVC
SVC()
```

```
In [16]: LR_acc = accuracy_score(y_test, LR.predict(X_test_vec))
NB_acc = accuracy_score(y_test, NB.predict(X_test_vec))
DT_acc = accuracy_score(y_test, DT.predict(X_test_vec))
RF_acc = accuracy_score(y_test, RF.predict(X_test_vec))
svc_acc = accuracy_score(y_test, svc.predict(X_test_vec))
```

```
In [17]: print("LR Accuracy: ", LR_acc)
print("NB Accuracy: ", NB_acc)
print("DT Accuracy: ", DT_acc)
print("RF Accuracy: ", RF_acc)
print("SVC Accuracy: ", svc_acc)
```

```
LR Accuracy: 0.967713004484305
NB Accuracy: 0.9766816143497757
DT Accuracy: 0.968609865470852
RF Accuracy: 0.97847533632287
SVC Accuracy: 0.9847533632286996
```

```
In [18]: LR_p = precision_score(y_test, LR.predict(X_test_vec))
        NB_p = precision_score(y_test, NB.predict(X_test_vec))
        DT_p = precision_score(y_test, DT.predict(X_test_vec))
        RF_p = precision_score(y_test, RF.predict(X_test_vec))
        svc_p = precision_score(y_test, svc.predict(X_test_vec))
```

```
In [19]: print("LR Precision: ", LR_p)
        print("NB Precision: ", NB_p)
        print("DT Precision: ", DT_p)
        print("RF Precision: ", RF_p)
        print("SVC Precision: ", svc_p)
```

```
LR Precision: 1.0
NB Precision: 1.0
DT Precision: 0.9014084507042254
RF Precision: 1.0
SVC Precision: 1.0
```


```
In [20]: LR_r = recall_score(y_test, LR.predict(X_test_vec))
        NB_r = recall_score(y_test, NB.predict(X_test_vec))
        DT_r = recall_score(y_test, DT.predict(X_test_vec))
        RF_r = recall_score(y_test, RF.predict(X_test_vec))
        svc_r = recall_score(y_test, svc.predict(X_test_vec))
```

```
In [21]: print("LR Recall: ", LR_r)
        print("NB Recall: ", NB_r)
        print("DT Recall: ", DT_r)
        print("RF Recall: ", RF_r)
        print("SVC Recall: ", svc_r)
```


```
LR Recall: 0.7583892617449665
NB Recall: 0.825503355704698
DT Recall: 0.8590604026845637
RF Recall: 0.8389261744966443
SVC Recall: 0.8859060402684564
```

```
In [22]: import seaborn as sns
        import matplotlib.pyplot as plt
```


```
In [23]: lr_cm = confusion_matrix(y_test, LR.predict(X_test_vec))
        plt.figure(figsize=(5, 4))
        sns.heatmap(lr_cm, annot=True, fmt='d', cmap='Blues',
                    xticklabels=['Not Spam', 'Spam'], yticklabels=['Not Spam', 'Spam'])
        plt.title('Confusion Matrix: Logistic Regression')
        plt.xlabel('Predicted')
        plt.ylabel('Actual')
        plt.show()
```


```
In [24]: nb_cm = confusion_matrix(y_test, NB.predict(X_test_vec))
plt.figure(figsize=(5, 4))
sns.heatmap(nb_cm, annot=True, fmt='d', cmap='Greens',
            xticklabels=['Not Spam', 'Spam'], yticklabels=['Not Spam', 'Spam'])
plt.title('Confusion Matrix: Naive Bayes')
plt.xlabel('Predicted')
plt.ylabel('Actual')
plt.show()
```


```
In [25]: dt_cm = confusion_matrix(y_test, DT.predict(X_test_vec))
plt.figure(figsize=(5, 4))
sns.heatmap(dt_cm, annot=True, fmt='d', cmap='Oranges',
            xticklabels=['Not Spam', 'Spam'], yticklabels=['Not Spam', 'Spam'])
plt.title('Confusion Matrix: Decision Tree')
plt.xlabel('Predicted')
plt.ylabel('Actual')
plt.show()
```



```
In [26]: rf_cm = confusion_matrix(y_test, RF.predict(X_test_vec))
plt.figure(figsize=(5, 4))
sns.heatmap(rf_cm, annot=True, fmt='d', cmap='Purples',
            xticklabels=['Not Spam', 'Spam'], yticklabels=['Not Spam', 'Spam'])
plt.title('Confusion Matrix: Random Forest')
plt.xlabel('Predicted')
plt.ylabel('Actual')
plt.show()
```



```
In [27]: svc_cm = confusion_matrix(y_test, svc.predict(X_test_vec))
plt.figure(figsize=(5, 4))
sns.heatmap(svc_cm, annot=True, fmt='d', cmap='Reds',
            xticklabels=['Not Spam', 'Spam'], yticklabels=['Not Spam', 'Spam'])
plt.title('Confusion Matrix: SVC')
plt.xlabel('Predicted')
plt.ylabel('Actual')
plt.show()
```



```
In [31]: import pickle


with open('/content/drive/MyDrive/AIML_lab_sem4/exp 5,6,7/exp5_models/logistic.pkl', 'wb') as f:
    pickle.dump(LR, f)

with open('/content/drive/MyDrive/AIML_lab_sem4/exp 5,6,7/exp5_models/naivebay.pkl', 'wb') as f:
    pickle.dump(NB, f)

with open('/content/drive/MyDrive/AIML_lab_sem4/exp 5,6,7/exp5_models/decision.pkl', 'wb') as f:
    pickle.dump(DT, f)

with open('/content/drive/MyDrive/AIML_lab_sem4/exp 5,6,7/exp5_models/randomforest.pkl', 'wb') as f:
    pickle.dump(RF, f)

with open('/content/drive/MyDrive/AIML_lab_sem4/exp 5,6,7/exp5_models/supportvect.pkl', 'wb') as f:
    pickle.dump(svc, f)
```



Experiment 6

Credit Risk Assessment

Aim : Build a credit scoring model to assess the creditworthiness of applicants using historical financial data and perform comparative analysis (Logistic Regression, Random Forest, XGBoost)

Problem Definition

Automated credit scoring is vital for financial institutions to mitigate potential losses and optimize lending decisions. Because the volume of applications makes manual review inefficient, we require robust classification models to predict applicant creditworthiness. This project aims to build a reliable credit filter by performing a comparative analysis of three algorithms: Logistic Regression, Random Forest, and Gradient Boosting. By evaluating these models, we seek to identify the architecture that most accurately categorizes "High," "Medium," and "Low" risk applicants, ensuring proactive risk management and improved financial stability.

Pipeline :

1. Problem Definition

The primary goal of our experiment is to develop a robust automated credit scoring model using advanced classification techniques:

- **Classification Task:** We must determine the creditworthiness of an applicant, categorizing them into risk levels such as High, Medium, or Low.
- **Objective:** We perform a detailed comparative analysis across three distinct algorithms Logistic Regression, Random Forest, and Gradient Boosting to identify the most reliable architecture for financial deployment.

2. Data Understanding

- **Dataset Content:** We utilize a collection of 500 samples of applicant profiles including historical financial health indicators.
- **Input Features (X):** We analyze raw data including age, annual income, employment years, credit score, loan amount, and late payments.
- **Target (y):** We use categorical labels representing "Credit Risk" (High, Medium, or Low).

3. Data Preprocessing

Since financial data requires specific cleaning, we put it through a rigorous transformation process:

- **Feature Selection:** We drop non-predictive identifiers, such as `Applicant_ID`, to focus on significant vocabulary and data.
- **Label Encoding:** We utilize a `LabelEncoder` to transform categorical risk labels into numerical vectors.
- **Standardization:** We apply a `StandardScaler` to normalize numerical features like income and loan amounts, preventing large values from dominating the model logic.

4. Train-Test Split

To ensure our filter works on applicants it has never seen before, we strategically divide the dataset:

- **Data Allocation:** We split the cleaned dataset into a Training set (80%) to build our models and a Test set (20%) to validate them.
- **Purpose:** We use this division to prevent "overfitting" and allow us to measure how accurately the filter will perform in a live financial environment.

5. Model Selection and Training

We initialize and train three different classification models on the same data to ensure a fair comparison:

- **Logistic Regression:** We use this as a baseline linear model to estimate risk probabilities based on weighted input features.
- **Random Forest:** We implement this ensemble method to construct multiple decision trees, improving stability and capturing non-linear relationships.
- **Gradient Boosting:** We task this model with correcting errors of previous trees sequentially to achieve high-fidelity predictions.

6. Predictions

- **Execution:** Once training is complete, we task each model with classifying the 100 unseen samples in our test set.
- **Output Generation:** We generate binary and multi-class prediction arrays for each algorithm to prepare for the evaluation phase.

7. Model evaluation and comparison :

The performance of each model is measured using Accuracy, Precision, and Recall. Based on the experimental outcomes.

	Accuracy	Precision	Recall	F1
Logistic regression	0.85	0.81	0.76	0.79
Random Forest	0.98	0.98	0.87	0.91
Gradient Boosting	1.00	1.00	1.00	1.00

We determine that Gradient Boosting is the best model among all tested as it scored the highest across all critical financial metrics and showed perfect classification capability on the test set.

8. Model Saving

The final step is ensuring our models are persistent for real-time filtering:

- We save all trained models, along with the scaler and label encoder, using the Pickle library to serialize the objects into permanent .pkl files.
- We generate these files to allow for immediate deployment in a banking system without the need for retraining.


```
In [ ]: import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import LabelEncoder, StandardScaler
from sklearn.metrics import accuracy_score, precision_score, recall_score, f1_
```

```
In [ ]: from sklearn.linear_model import LogisticRegression
from sklearn.ensemble import RandomForestClassifier
from sklearn.ensemble import GradientBoostingClassifier
```

```
In [ ]: df = pd.read_csv('/content/drive/MyDrive/AI ML_lab_sem4/exp 5,6,7/credit_risk_4')
df.head()
```

```
Out[ ]:   Applicant_ID  Age  Annual_Income  Employment_Years  Credit_Score  Loan_Amount
0             1    59         153267             28           818           15000
1             2    49         232745             0           436           15000
2             3    35         974945             19           797           60000
3             4    63         307164             29           776           80000
4             5    28         685626             6           893           40000
```

```
In [ ]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 500 entries, 0 to 499
Data columns (total 11 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   Applicant_ID                          500 non-null    int64
1   Age                                    500 non-null    int64
2   Annual_Income                         500 non-null    int64
3   Employment_Years                      500 non-null    int64
4   Credit_Score                          500 non-null    int64
5   Loan_Amount                           500 non-null    int64
6   Loan_Term_Months                      500 non-null    int64
7   Existing_Loans_Count                  500 non-null    int64
8   Debt_to_Income_Ratio                  500 non-null    float64
9   Late_Payments_Last_2Yrs              500 non-null    int64
10  Credit_Risk                           500 non-null    object
dtypes: float64(1), int64(9), object(1)
memory usage: 43.1+ KB
```

```
In [ ]: print(df['Credit_Risk'].value_counts())
```

```
Credit_Risk
High      315
Medium    164
Low        21
Name: count, dtype: int64
```

```
In [ ]: X = df.drop(['Applicant_ID', 'Credit_Risk'], axis=1)
        y = df['Credit_Risk']
```

```
In [ ]: le = LabelEncoder()
        y = le.fit_transform(y)
```

```
In [ ]: X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)
```

```
In [ ]: scaler = StandardScaler()
        X_train_scaled = scaler.fit_transform(X_train)
        X_test_scaled = scaler.transform(X_test)
```

```
In [ ]: LR = LogisticRegression(max_iter=1000)
        LR.fit(X_train_scaled, y_train)
```

```
Out[ ]: *      LogisticRegression
        LogisticRegression(max_iter=1000)
```

```
In [ ]: RF = RandomForestClassifier(random_state=42)
        RF.fit(X_train, y_train)
```

```
Out[ ]: *      RandomForestClassifier
        RandomForestClassifier(random_state=42)
```

```
In [ ]: GB = GradientBoostingClassifier(random_state=42)
        GB.fit(X_train, y_train)
```

```
Out[ ]: *      GradientBoostingClassifier
        GradientBoostingClassifier(random_state=42)
```

```
In [ ]: LR_acc = accuracy_score(y_test, LR.predict(X_test_scaled))
        RF_acc = accuracy_score(y_test, RF.predict(X_test))
        GB_acc = accuracy_score(y_test, GB.predict(X_test))
```

```
In [ ]: print(f"Logistic Regression Accuracy: {LR_acc:.4f}")
        print(f"Random Forest Accuracy: {RF_acc:.4f}")
        print(f"Gradient Boosting Accuracy: {GB_acc:.4f}")
```

```
Logistic Regression Accuracy: 0.8500
Random Forest Accuracy: 0.9800
Gradient Boosting Accuracy: 1.0000
```

```
In [ ]: LR_p = precision_score(y_test, LR.predict(X_test_scaled), average='macro')
        RF_p = precision_score(y_test, RF.predict(X_test), average='macro')
        GB_p = precision_score(y_test, GB.predict(X_test), average='macro')
```

```

LR_r = recall_score(y_test, LR.predict(X_test_scaled), average='macro')
RF_r = recall_score(y_test, RF.predict(X_test), average='macro')
GB_r = recall_score(y_test, GB.predict(X_test), average='macro')

LR_f1 = f1_score(y_test, LR.predict(X_test_scaled), average='macro')
RF_f1 = f1_score(y_test, RF.predict(X_test), average='macro')
GB_f1 = f1_score(y_test, GB.predict(X_test), average='macro')

```

```

In [ ]: print("\n--- Detailed Metrics (Precision, Recall, F1) ---")
print(f"Logistic Regression -> Precision: {LR_p:.2f}, Recall: {LR_r:.2f}, F1: {LR_f1:.2f}")
print(f"Random Forest -> Precision: {RF_p:.2f}, Recall: {RF_r:.2f}, F1: {RF_f1:.2f}")
print(f"Gradient Boosting -> Precision: {GB_p:.2f}, Recall: {GB_r:.2f}, F1: {GB_f1:.2f}")

--- Detailed Metrics (Precision, Recall, F1) ---
LR -> Precision: 0.81, Recall: 0.76, F1: 0.79
RF -> Precision: 0.98, Recall: 0.87, F1: 0.91
GB -> Precision: 1.00, Recall: 1.00, F1: 1.00


```

```

In [ ]: lr_cm = confusion_matrix(y_test, LR.predict(X_test_scaled))


plt.figure(figsize=(6, 5))
sns.heatmap(lr_cm, annot=True, fmt='d', cmap='Blues',
            xticklabels=le.classes_, yticklabels=le.classes_)
plt.title('Confusion Matrix: Logistic Regression')
plt.xlabel('Predicted Label')
plt.ylabel('True Label')
plt.show()

```




```
In [ ]: rf_cm = confusion_matrix(y_test, RF.predict(X_test))

plt.figure(figsize=(6, 5))
sns.heatmap(rf_cm, annot=True, fmt='d', cmap='Greens',
            xticklabels=le.classes_, yticklabels=le.classes_)
plt.title('Confusion Matrix: Random Forest')
plt.xlabel('Predicted Label')
plt.ylabel('True Label')
plt.show()
```



```
In [ ]: gb_cm = confusion_matrix(y_test, GB.predict(X_test))

plt.figure(figsize=(6, 5))
sns.heatmap(gb_cm, annot=True, fmt='d', cmap='Purples',
            xticklabels=le.classes_, yticklabels=le.classes_)
plt.title('Confusion Matrix: Gradient Boosting')
plt.xlabel('Predicted Label')
plt.ylabel('True Label')
plt.show()
```

```
In [ ]: import pickle





with open('/content/drive/MyDrive/AIML_lab_sen4/exp 5,6,7/exp6_models/lr_model.pkl', 'wb') as f:
    pickle.dump(LR, f)

with open('/content/drive/MyDrive/AIML_lab_sen4/exp 5,6,7/exp6_models/scaler.pkl', 'wb') as f:
    pickle.dump(scaler, f)

with open('/content/drive/MyDrive/AIML_lab_sen4/exp 5,6,7/exp6_models/rf_model.pkl', 'wb') as f:
    pickle.dump(RF, f)

with open('/content/drive/MyDrive/AIML_lab_sen4/exp 5,6,7/exp6_models/gb_model.pkl', 'wb') as f:
    pickle.dump(GB, f)

with open('/content/drive/MyDrive/AIML_lab_sen4/exp 5,6,7/exp6_models/label_encoder.pkl', 'wb') as f:
    pickle.dump(le, f)
```



Experiment 7

Anomaly Detection

Aim : Implement an anomaly detection system for detecting outliers in data (e.g., fraud detection) and perform comparative analysis. (Isolation Forest, Local Outlier Factor, One-Class SVM)

Problem :

Anomaly detection is critical for identifying rare items, events, or observations which raise suspicions by differing significantly from the majority of the data. Because the volume of transactions makes manual inspection for fraud or errors impossible, we require automated unsupervised models to detect outliers. This project aims to build a robust detection system by performing a comparative analysis of three algorithms: Isolation Forest, Local Outlier Factor (LOF), and One-Class SVM. By evaluating these models, we seek to identify the architecture that most accurately distinguishes between "Normal" data and "Anomalies," ensuring proactive security and data integrity.

Pipeline :

1. Problem Definition

The primary goal of our experiment is to develop a robust automated anomaly detection system using advanced unsupervised learning techniques:

- **Classification Task:** We must determine if a data point is Normal (1) or an Anomaly (-1).
- **Objective:** We perform a detailed comparative analysis across three distinct algorithms—Isolation Forest, Local Outlier Factor, and One-Class SVM to identify the most reliable architecture for outlier detection.

2. Data Understanding

- **Dataset Content:** We utilize a collection of 500 samples consisting of features like Transaction Amount, Frequency, and Time.
- **Input Features (X):** We analyze numerical data representing behavioral patterns (e.g., Feature_1, Feature_2).
- **Target (y):** We use labels where 1 represents a "Normal" observation and -1 represents an "Anomaly" (Outlier).

3. Data Preprocessing

Since anomaly detection models are sensitive to the scale of data, we put it through a rigorous transformation process:

- **Cleaning:** We remove non-predictive columns and ensure the data is formatted for unsupervised learning.
- **Standardization:** We apply a `StandardScaler` to normalize numerical features, ensuring that features with larger ranges (like transaction amounts) do not disproportionately influence the distance-based models like LOF and SVM.

4. Train-Test Split

To evaluate how well the system identifies new outliers, we strategically divide the dataset:

- **Data Allocation:** We split the dataset into a Training set (80%) to establish the "normal" baseline and a Test set (20%) to validate detection accuracy.
- **Purpose:** This division allows us to measure the model's ability to generalize and identify anomalies in unseen data.

5. Model Selection and Training

We initialize and train three different anomaly detection models on the same data to ensure a fair comparison:

- **Isolation Forest:** We use this tree-based model that explicitly isolates anomalies by randomly selecting a feature and a split value.
- **Local Outlier Factor (LOF):** We implement this density-based method that identifies outliers by comparing the local density of a point to its neighbors.
- **One-Class SVM:** We task this model with learning a decision boundary that encompasses the "normal" data points in a high-dimensional space.

6. Predictions

- Execution: Once training is complete, we task each model with classifying the 100 unseen samples in our test set.
- Output Generation: We generate prediction arrays where each algorithm flags points as either 1 (Normal) or -1 (Anomaly).

7. Model evaluation and comparison :

The performance of each model is measured using Accuracy, Precision, Recall, and F1-Score. Based on the experimental outcomes:

	Accuracy	Precision	Recall	F1
Isolation forest	0.96	0.7	0.7	0.7
Local Outlier Factor	0.88	0	0	0
One Class SVM	0.93	0.41	0.43	0.42

We determine that Isolation Forest is the best model among all tested as it achieved the highest accuracy and balanced precision and recall, making it the most effective at isolating outliers without misclassifying normal data.

8. Model Saving

The final step is ensuring our models are persistent for real-time monitoring:

- Implementation: We save all trained models, along with the anomaly scaler, using the Pickle library to serialize the objects into permanent .pkl files.
- Final Output: We generate these files to allow for immediate deployment in a fraud detection system for real-time filtering without the need for retraining.

```
In [ ]: import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt
import pickle
from sklearn.preprocessing import StandardScaler
from sklearn.ensemble import IsolationForest
from sklearn.neighbors import LocalOutlierFactor
from sklearn.svm import OneClassSVM
from sklearn.metrics import accuracy_score, precision_score, recall_score, fl_
```

```
In [ ]: df = pd.read_csv('/content/drive/MyDrive/AIML_lab_sem4/exp 5,6,7/anomaly_detect.csv')
```

```
In [ ]: df.head()
```

```
Out[ ]:   Feature_1  Feature_2  Anomaly_Label
0   43.259073   53.716320             0
1   43.815246   43.397717             0
2   46.942411   42.966695             0
3   48.911594   55.493884             0
4   52.575238   69.263657             0
```

```
In [ ]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 500 entries, 0 to 499
Data columns (total 3 columns):
#   Column          Non-Null Count  Dtype
---  -
0   Feature_1       500 non-null   float64
1   Feature_2       500 non-null   float64
2   Anomaly_Label   500 non-null   int64
dtypes: float64(2), int64(1)
memory usage: 11.8 KB
```

```
In [ ]: df['Anomaly_Label'].value_counts()
```

```
Out[ ]:      count
Anomaly_Label
0         470
1          30
```

dtype: int64

```
In [ ]: X = df[['Feature_1', 'Feature_2']]
```

```

In [ ]: y_true = df['Anomaly_Label']

In [ ]: scaler = StandardScaler()
X_scaled = scaler.fit_transform(X)

In [ ]: ISO = IsolationForest(contamination=0.06, random_state=42)
ISO_preds_raw = ISO.fit_predict(X_scaled)
ISO_preds = np.where(ISO_preds_raw == 1, 0, 1)

In [ ]: LOF = LocalOutlierFactor(contamination=0.06)
LOF_preds_raw = LOF.fit_predict(X_scaled)
LOF_preds = np.where(LOF_preds_raw == 1, 0, 1)

In [ ]: SVM = OneClassSVM(nu=0.06, kernel="rbf", gamma=0.1)
SVM_preds_raw = SVM.fit_predict(X_scaled)
SVM_preds = np.where(SVM_preds_raw == 1, 0, 1)

In [ ]: ISO_acc = accuracy_score(y_true, ISO_preds)
LOF_acc = accuracy_score(y_true, LOF_preds)
SVM_acc = accuracy_score(y_true, SVM_preds)

In [ ]: print("ISO Accuracy: ", ISO_acc)
print("LOF Accuracy: ", LOF_acc)
print("SVM Accuracy: ", SVM_acc)

ISO Accuracy:  0.964
LOF Accuracy:  0.88
SVM Accuracy:  0.93

In [ ]: ISO_p = precision_score(y_true, ISO_preds)
LOF_p = precision_score(y_true, LOF_preds)
SVM_p = precision_score(y_true, SVM_preds)

In [ ]: print("ISO Precision: ", ISO_p)
print("LOF Precision: ", LOF_p)
print("SVM Precision: ", SVM_p)

ISO Precision:  0.7
LOF Precision:  0.0
SVM Precision:  0.41935483870967744

In [ ]: ISO_r = recall_score(y_true, ISO_preds)
LOF_r = recall_score(y_true, LOF_preds)
SVM_r = recall_score(y_true, SVM_preds)

In [ ]: print("ISO Recall: ", ISO_r)
print("LOF Recall: ", LOF_r)
print("SVM Recall: ", SVM_r)

ISO Recall:  0.7
LOF Recall:  0.0
SVM Recall:  0.43333333333333335

```

```
In [ ]: ISO_f1 = f1_score(y_true, ISO_preds)
        LOF_f1 = f1_score(y_true, LOF_preds)
        SVM_f1 = f1_score(y_true, SVM_preds)
```


```
In [ ]: print("ISO F1: ", ISO_f1)
        print("LOF F1: ", LOF_f1)
        print("SVM F1: ", SVM_f1)
```

```
ISO F1:  0.7
LOF F1:  0.0
SVM F1:  0.4262295081967213
```

```
In [ ]: from sklearn.metrics import confusion_matrix
        import seaborn as sns
        import matplotlib.pyplot as plt
```


```
In [ ]: iso_cm = confusion_matrix(y_true, ISO_preds)
        plt.figure(figsize=(6, 5))
        sns.heatmap(iso_cm, annot=True, fmt='d', cmap='Blues',
                    xticklabels=['Normal', 'Anomaly'],
                    yticklabels=['Normal', 'Anomaly'])

        plt.title('Confusion Matrix: Isolation Forest')
        plt.xlabel('Predicted Label')
        plt.ylabel('True Label')
        plt.show()
```



```
In [ ]: lof_cm = confusion_matrix(y_true, LOF_preds)
plt.figure(figsize=(6, 5))
sns.heatmap(lof_cm, annot=True, fmt='d', cmap='Greens',
            xticklabels=['Normal', 'Anomaly'],
            yticklabels=['Normal', 'Anomaly'])


plt.title('Confusion Matrix: Local Outlier Factor')
plt.xlabel('Predicted Label')
plt.ylabel('True Label')
plt.show()
```



```
In [ ]: svm_cm = confusion_matrix(y_true, SVM_preds)


plt.figure(figsize=(6, 5))
sns.heatmap(svm_cm, annot=True, fmt='d', cmap='Purples',
            xticklabels=['Normal', 'Anomaly'],
            yticklabels=['Normal', 'Anomaly'])

plt.title('Confusion Matrix: One-Class SVM')
plt.xlabel('Predicted Label')
plt.ylabel('True Label')
plt.show()
```



```
In [ ]: results_df = pd.DataFrame({
    'Model': ['Isolation Forest', 'LOF', 'One-Class SVM'],
    'Accuracy': [ISO_acc, LOF_acc, SVM_acc],
    'Precision': [ISO_p, LOF_p, SVM_p],
    'Recall': [ISO_r, LOF_r, SVM_r],
    'F1 Score': [ISO_f1, LOF_f1, SVM_f1]
})
df_plot = results_df.melt(id_vars='Model', var_name='Metric', value_name='Score')

plt.figure(figsize=(10, 6))
sns.barplot(x='Model', y='Score', hue='Metric', data=df_plot)
plt.title('Performance Comparison: Anomaly Detection')
plt.ylim(0, 1.1)
plt.show()
```

```
In [ ]: with open('/content/drive/MyDrive/AIML_lab_sen4/exp 5,6,7/Exp7_models/iso_fore
        pickle.dump(Iso, f)

        with open('/content/drive/MyDrive/AIML_lab_sen4/exp 5,6,7/Exp7_models/lof_model
        pickle.dump(LOF, f)

        with open('/content/drive/MyDrive/AIML_lab_sen4/exp 5,6,7/Exp7_models/oc_svm.p
        pickle.dump(SVM, f)

        with open('/content/drive/MyDrive/AIML_lab_sen4/exp 5,6,7/Exp7_models/anomaly_
        pickle.dump(scaler, f)

        print("\nAll models and the scaler have been saved successfully.")
```

All models and the scaler have been saved successfully.

Experiment 8

Student Performance Level Analysis

Aim : Implement Multiclass Classification models for students Performance Level analysis and perform comparative analysis. (Random Forest, Decision Tree, Multinomial Logistic Regression, XGBoost, K-Nearest Neighbors).

1. Problem Definition

The objective of this experiment is to develop an automated system to predict a student's **Performance Level** based on various academic and behavioral metrics.

- **Classification Task:** This is a multiclass classification problem where students are categorized into distinct levels (e.g., 0, 1, or 2).
- **Objective:** Perform a comparative analysis across five distinct algorithms Multinomial Logistic Regression, Decision Tree, Random Forest, Gradient Boosting (XGBoost/GBM), and K-Nearest Neighbors (KNN) to determine which model best predicts academic outcomes.

2. Data Understanding

- **Dataset Content:** A collection of 500 student samples with academic indicators.
- **Input Features (X):**
 - *Numerical:* Study Hours, Attendance Percentage, Assignment Score, Internal Marks.
 - *Categorical:* Participation (Low/Medium/High), Internet Access (Yes/No), Previous Grade (A/B/C).
- **Target (y):** Performance_Level (Multiclass labels representing different tiers of academic achievement).

3. Data Preprocessing

To prepare the dataset for multiclass classification, a rigorous transformation process was applied to ensure the data is in a format compatible with both linear and distance-based algorithms:

- **Feature-Target Separation:** After encoding, the dataset was divided into:
 - Input Features (X): All academic and behavioral features excluding the target.

- Target (y): The Performance_Level label.
- **Standardization:** We applied StandardScaler to normalize the numerical features.
- **Label encoding :** We perform label encoding on the categorical data to convert it into numerical data.

4. Train-Test Split

- **Data Allocation:** The dataset was partitioned into a Training set (80%) to train the models and a Test set (20%) to evaluate their performance on unseen data.
- **Random State:** A fixed random_state=42 was used to ensure the results are reproducible across different runs.

5. Model Selection and Training

Five different classification architectures were initialized and trained:

1. **Multinomial Logistic Regression:** A linear model adapted for multiclass settings using the softmax function.
2. **Decision Tree:** A non-linear model that splits data based on feature thresholds to create a tree-like decision structure.
3. **Random Forest:** An ensemble of decision trees that reduces overfitting by averaging multiple "votes."
4. **Gradient Boosting:** An iterative ensemble technique that builds new trees to correct the errors made by previous ones.
5. **K-Nearest Neighbors (KNN):** A distance-based model that classifies points based on the majority label of their nearest neighbors.

6. Predictions

- Testing data is used to do prediction in the trained model.
- The predicted output is stored to further use than in model evaluation.

7. Model Evaluation and Comparison

The models were evaluated using **Accuracy**, **Precision**, **Recall**, and **F1-Score** (using 'macro' averaging to account for class distribution) / as well as AUC ROC curve.

	Accuracy	Precision	Recall	F1-Score
Logistic Regression	0.9700	0.6474	0.6599	0.6533
Decision Tree	0.9000	0.9320	0.9320	0.9320
Random Forest	0.9400	0.6269	0.6395	0.6331
Gradient Boosting	0.9500	0.9661	0.9660	0.9660
K-Nearest Neighbors	0.7900	0.5272	0.5374	0.5321

Roc-Auc curve analysis

	Class 0	Class1	Class2
Logistic Regression	1.00	0.96	0.86
Decision Tree	0.9	0.90	1.00
Random Forest	1.00	0.98	1.00
Gradient Boosting	0.99	0.99	1.00
K-Nearest Neighbors	0.92	0.89	0.50

Conclusion: While Logistic Regression achieved the highest raw accuracy, **Gradient Boosting** is the most robust model for this dataset, as it provides the best balance between high accuracy and superior Precision/Recall/F1-Scores.

8. Model Saving

To make the system deployable for real-time student analysis:

- Implementation: All five trained models and the StandardScaler were serialized into .pkl files using the Pickle library.
- Final Output: These files allow the prediction system to be reloaded in a production environment (like a school dashboard) without retraining the models.

```
In [65]: import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt
import pickle
from sklearn.metrics import classification_report
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import LabelEncoder, StandardScaler
from sklearn.metrics import accuracy_score, precision_score, recall_score, f1_
```

```
In [66]: from sklearn.linear_model import LogisticRegression
from sklearn.tree import DecisionTreeClassifier
from sklearn.ensemble import RandomForestClassifier, GradientBoostingClassifier
from sklearn.neighbors import KNeighborsClassifier
```

```
In [67]: df = pd.read_csv('/content/drive/MyDrive/AIML_lab_sem4/exp 5,6,7/multiclass_cl
```

```
In [68]: df.head()
```

```
Out[68]:
```

	Study_Hours	Attendance_Percentage	Assignment_Score	Internal_Marks	Par
0	22.48	84.26	86.79	76.68	
1	19.31	94.09	81.10	56.73	
2	23.24	61.01	70.72	52.73	
3	27.62	80.63	62.24	64.95	
4	18.83	68.49	78.38	62.45	

```
In [69]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 500 entries, 0 to 499
Data columns (total 8 columns):
#   Column                Non-Null Count  Dtype
---  -
0   Study_Hours           500 non-null   float64
1   Attendance_Percentage 500 non-null   float64
2   Assignment_Score      500 non-null   float64
3   Internal_Marks        500 non-null   float64
4   Participation          500 non-null   object
5   Internet_Access       500 non-null   object
6   Previous_Grade        500 non-null   object
7   Performance_Level     500 non-null   int64
dtypes: float64(4), int64(1), object(3)
memory usage: 31.4+ KB
```

```
In [70]: le = LabelEncoder()
df['Participation'] = le.fit_transform(df['Participation'])
df['Internet_Access'] = le.fit_transform(df['Internet_Access'])
df['Previous_Grade'] = le.fit_transform(df['Previous_Grade'])
```

```
In [71]: df.head()
```

```
Out[71]:
```

	Study_Hours	Attendance_Percentage	Assignment_Score	Internal_Marks	Par
0	22.48	84.26	86.79	76.68	
1	19.31	94.09	81.10	56.73	
2	23.24	61.01	70.72	52.73	
3	27.62	80.63	62.24	64.95	
4	18.83	68.49	78.38	62.45	

```
In [72]: X = df.drop('Performance_Level', axis=1)
y = df['Performance_Level']
```

```
In [73]: X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)
```

```
In [74]: scaler = StandardScaler()
X_train_scaled = scaler.fit_transform(X_train)
X_test_scaled = scaler.transform(X_test)
```

```
In [75]: lr = LogisticRegression(multi_class='multinomial', max_iter=1000)
lr.fit(X_train_scaled, y_train)
lr_pred = lr.predict(X_test_scaled)
```

```
/usr/local/lib/python3.12/dist-packages/sklearn/linear_model/_logistic.py:1247:
FutureWarning: 'multi_class' was deprecated in version 1.5 and will be removed
in 1.7. From then on, it will always use 'multinomial'. Leave it to its default
value to avoid this warning.
  warnings.warn()
```

```
In [76]: print(f"Accuracy: {accuracy_score(y_test, lr_pred):.4f}")
print(f"Precision: {precision_score(y_test, lr_pred, average='macro', zero_division=0):.4f}")
print(f"Recall: {recall_score(y_test, lr_pred, average='macro'): .4f}")
print(f"F1-Score: {f1_score(y_test, lr_pred, average='macro'): .4f}")
```

```
Accuracy: 0.9700
Precision: 0.6474
Recall: 0.6599
F1-Score: 0.6533
```

```
In [77]: print('Classification Report for Logistic Regression:')
print(classification_report(y_test, lr_pred, zero_division=0))
```

```

Classification Report for Logistic Regression:
      precision    recall  f1-score   support

     0       1.00      0.98      0.99         49
     1       0.94      1.00      0.97         49
     2       0.00      0.00      0.00          2

 accuracy      0.97         100
 macro avg      0.65         100
 weighted avg    0.95         100

```

```

In [78]: dt = DecisionTreeClassifier(random_state=42)
dt.fit(X_train, y_train)
dt_pred = dt.predict(X_test)

```

```

In [79]: print(f"Accuracy: {accuracy_score(y_test, dt_pred):.4f}")
print(f"Precision: {precision_score(y_test, dt_pred, average='macro', zero_div
print(f"Recall: {recall_score(y_test, dt_pred, average='macro'):.4f}")
print(f"F1-Score: {f1_score(y_test, dt_pred, average='macro'):.4f}")

```

```

Accuracy: 0.9000
Precision: 0.9320
Recall: 0.9320
F1-Score: 0.9320

```

```

In [80]: print('Classification Report for Decision Tree:')
print(classification_report(y_test, dt_pred, zero_division=0))

```

```

Classification Report for Decision Tree:
      precision    recall  f1-score   support

     0       0.90      0.90      0.90         49
     1       0.90      0.90      0.90         49
     2       1.00      1.00      1.00          2

 accuracy      0.90         100
 macro avg      0.93         100
 weighted avg    0.90         100

```

```

In [81]: rf = RandomForestClassifier(random_state=42)
rf.fit(X_train, y_train)
rf_pred = rf.predict(X_test)

```

```

In [82]: print(f"Accuracy: {accuracy_score(y_test, rf_pred):.4f}")
print(f"Precision: {precision_score(y_test, rf_pred, average='macro', zero_div
print(f"Recall: {recall_score(y_test, rf_pred, average='macro'):.4f}")
print(f"F1-Score: {f1_score(y_test, rf_pred, average='macro'):.4f}")

```

```

Accuracy: 0.9400
Precision: 0.6269
Recall: 0.6395
F1-Score: 0.6331

```



```
In [83]: print('Classification Report for Random Forest:')
print(classification_report(y_test, rf_pred, zero_division=0))
```

```
Classification Report for Random Forest:
              precision    recall  f1-score   support

    0           0.96       0.96      0.96         49
    1           0.92       0.96      0.94         49
    2           0.00       0.00      0.00          2

 accuracy          0.94         100
 macro avg         0.63         0.64         0.63         100
 weighted avg      0.92         0.94         0.93         100
```

```
In [84]: gb = GradientBoostingClassifier(random_state=42)
gb.fit(X_train, y_train)
gb_pred = gb.predict(X_test)
```

```
In [85]: print(f"Accuracy: {accuracy_score(y_test, gb_pred):.4f}")
print(f"Precision: {precision_score(y_test, gb_pred, average='macro', zero_division=0):.4f}")
print(f"Recall: {recall_score(y_test, gb_pred, average='macro'):~.4f}")
print(f"F1-Score: {f1_score(y_test, gb_pred, average='macro'):~.4f}")
```

```
Accuracy: 0.9500
Precision: 0.9661
Recall: 0.9660
F1-Score: 0.9660
```

```
In [86]: print('Classification Report for Gradient Boosting:')
print(classification_report(y_test, gb_pred, zero_division=0))
```

```
Classification Report for Gradient Boosting:
              precision    recall  f1-score   support

    0           0.94       0.96      0.95         49
    1           0.96       0.94      0.95         49
    2           1.00       1.00      1.00          2

 accuracy          0.95         100
 macro avg         0.97         0.97         0.97         100
 weighted avg      0.95         0.95         0.95         100
```

```
In [87]: knn = KNeighborsClassifier()
knn.fit(X_train_scaled, y_train)
knn_pred = knn.predict(X_test_scaled)
```

```
In [88]: print(f"Accuracy: {accuracy_score(y_test, knn_pred):.4f}")
print(f"Precision: {precision_score(y_test, knn_pred, average='macro', zero_division=0):.4f}")
print(f"Recall: {recall_score(y_test, knn_pred, average='macro'):~.4f}")
print(f"F1-Score: {f1_score(y_test, knn_pred, average='macro'):~.4f}")
```


Accuracy: 0.7900
Precision: 0.5272
Recall: 0.5374
F1-Score: 0.5321


```
In [89]: print('Classification Report for K-Nearest Neighbors:')  
print(classification_report(y_test, knn_pred, zero_division=0))
```

Classification Report for K-Nearest Neighbors:

	precision	recall	f1-score	support
0	0.81	0.80	0.80	49
1	0.77	0.82	0.79	49
2	0.00	0.00	0.00	2
accuracy			0.79	100
macro avg	0.53	0.54	0.53	100
weighted avg	0.78	0.79	0.78	100

```
In [90]: from sklearn.metrics import confusion_matrix, ConfusionMatrixDisplay  
import matplotlib.pyplot as plt  
  
LR_pred = lr_pred  
  
print(f'Accuracy: {accuracy_score(y_test, LR_pred):.4f}')  
print(f'Precision: {precision_score(y_test, LR_pred, average='macro', zero_division=0):.4f}')  
print(f'Recall: {recall_score(y_test, LR_pred, average='macro'): .4f}')  
print(f'F1-Score: {f1_score(y_test, LR_pred, average='macro'): .4f}')  
  
cm_lr = confusion_matrix(y_test, LR_pred)  
disp_lr = ConfusionMatrixDisplay(confusion_matrix=cm_lr)  
disp_lr.plot(cmap='Blues')  
plt.title('Confusion Matrix: Logistic Regression')  
plt.show()
```

Accuracy: 0.9700
Precision: 0.6474
Recall: 0.6599
F1-Score: 0.6533




```
In [91]: y_pred_proba_lr = lr.predict_proba(X_test_scaled)

fig_lr, ax_lr = plt.subplots(figsize=(8, 6))

for i in range(n_classes):
    RocCurveDisplay.from_predictions(
        y_test_binarized[:, i],
        y_pred_proba_lr[:, i],
        name=f'ROC curve for class {classes[i]}',
        ax=ax_lr,
    )

ax_lr.plot([0, 1], [0, 1], linestyle='--', lw=2, color='r', label='Chance Level')
ax_lr.set_title('ROC AUC Curve: Logistic Regression (One-vs-Rest)')
ax_lr.set_xlabel('False Positive Rate')
ax_lr.set_ylabel('True Positive Rate')
ax_lr.grid(linestyle='--', alpha=0.7)
ax_lr.legend(loc='lower right')
plt.show()
```




In [92]: DT_pred = dt_pred

```
print(f"Accuracy: {accuracy_score(y_test, DT_pred):.4f}")
print(f"Precision: {precision_score(y_test, DT_pred, average='macro', zero_div=0):.4f}")
print(f"Recall: {recall_score(y_test, DT_pred, average='macro'):~.4f}")
print(f"F1-Score: {f1_score(y_test, DT_pred, average='macro'):~.4f}")

cm_dt = confusion_matrix(y_test, DT_pred)
disp_dt = ConfusionMatrixDisplay(confusion_matrix=cm_dt)
disp_dt.plot(cmap='Greens')
plt.title('Confusion Matrix: Decision Tree')
plt.show()
```

```
Accuracy: 0.9000
Precision: 0.9320
Recall: 0.9320
F1-Score: 0.9320
```




```
In [93]: y_pred_proba_dt = dt.predict_proba(X_test)

fig_dt, ax_dt = plt.subplots(figsize=(8, 6))

for i in range(n_classes):
    RocCurveDisplay.from_predictions(
        y_test_binarized[:, i],
        y_pred_proba_dt[:, i],
        name=f'ROC curve for class {classes[i]}',
        ax=ax_dt,
    )

ax_dt.plot([0, 1], [0, 1], linestyle='--', lw=2, color='r', label='Chance Level')
ax_dt.set_title('ROC AUC Curve: Decision Tree (One-vs-Rest)')
ax_dt.set_xlabel('False Positive Rate')
ax_dt.set_ylabel('True Positive Rate')
ax_dt.grid(linestyle='--', alpha=0.7)
ax_dt.legend(loc='lower right')
plt.show()
```




```
In [94]: RF_pred = rf_pred
```

```
print(f"Accuracy: {accuracy_score(y_test, RF_pred):.4f}")
print(f"Precision: {precision_score(y_test, RF_pred, average='macro', zero_div=0):.4f}")
print(f"Recall: {recall_score(y_test, RF_pred, average='macro'): .4f}")
print(f"F1-Score: {f1_score(y_test, RF_pred, average='macro'): .4f}")

cm_rf = confusion_matrix(y_test, RF_pred)
disp_rf = ConfusionMatrixDisplay(confusion_matrix=cm_rf)
disp_rf.plot(cmap='Oranges')
plt.title('Confusion Matrix: Random Forest')
plt.show()
```

```
Accuracy: 0.9400
Precision: 0.6269
Recall: 0.6395
F1-Score: 0.6331
```




```
In [95]: y_pred_proba_rf = rf.predict_proba(X_test)

fig_rf, ax_rf = plt.subplots(figsize=(8, 6))

for i in range(n_classes):
    RocCurveDisplay.from_predictions(
        y_test_binarized[:, i],
        y_pred_proba_rf[:, i],
        name=f'ROC curve for class {classes[i]}',
        ax=ax_rf,
    )

ax_rf.plot([0, 1], [0, 1], linestyle='--', lw=2, color='r', label='Chance Level')
ax_rf.set_title('ROC AUC Curve: Random Forest (One-vs-Rest)')
ax_rf.set_xlabel('False Positive Rate')
ax_rf.set_ylabel('True Positive Rate')
ax_rf.grid(linestyle='--', alpha=0.7)
ax_rf.legend(loc='lower right')
plt.show()
```





```
In [96]: GB_pred = gb_pred
```

```
print(f"Accuracy: {accuracy_score(y_test, GB_pred):.4f}")
print(f"Precision: {precision_score(y_test, GB_pred, average='macro', zero_div
print(f"Recall: {recall_score(y_test, GB_pred, average='macro'):.4f}")
print(f"F1-Score: {f1_score(y_test, GB_pred, average='macro'):.4f}")

cm_gb = confusion_matrix(y_test, GB_pred)
disp_gb = ConfusionMatrixDisplay(confusion_matrix=cm_gb)
disp_gb.plot(cmap='Reds')
plt.title('Confusion Matrix: Gradient Boosting')
plt.show()
```

```
Accuracy: 0.9500
Precision: 0.9661
Recall: 0.9660
F1-Score: 0.9660
```




```
In [97]: y_pred_proba_gb = gb.predict_proba(X_test)

fig_gb, ax_gb = plt.subplots(figsize=(8, 6))

for i in range(n_classes):
    RocCurveDisplay.from_predictions(
        y_test_binarized[:, i],
        y_pred_proba_gb[:, i],
        name=f'ROC curve for class {classes[i]}',
        ax=ax_gb,
    )

ax_gb.plot([0, 1], [0, 1], linestyle='--', lw=2, color='r', label='Chance Level')
ax_gb.set_title('ROC AUC Curve: Gradient Boosting (One-vs-Rest)')
ax_gb.set_xlabel('False Positive Rate')
ax_gb.set_ylabel('True Positive Rate')
ax_gb.grid(linestyle='--', alpha=0.7)
ax_gb.legend(loc='lower right')
plt.show()
```





In [96]: `KNN_pred = knn_pred`

```
print(f"Accuracy: {accuracy_score(y_test, KNN_pred):.4f}")
print(f"Precision: {precision_score(y_test, KNN_pred, average='macro', zero_division=0):.4f}")
print(f"Recall: {recall_score(y_test, KNN_pred, average='macro'):.4f}")
print(f"F1-Score: {f1_score(y_test, KNN_pred, average='macro'):.4f}")

cm_knn = confusion_matrix(y_test, KNN_pred)
disp_knn = ConfusionMatrixDisplay(confusion_matrix=cm_knn)
disp_knn.plot(cmap='Purples')
plt.title('Confusion Matrix: KNN')
plt.show()
```

```
Accuracy: 0.7900
Precision: 0.5272
Recall: 0.5374
F1-Score: 0.5321
```




```
In [99]: from sklearn.metrics import RocCurveDisplay
from sklearn.preprocessing import label_binarize
import numpy as np
import matplotlib.pyplot as plt
y_pred_proba_knn = knn.predict_proba(X_test_scaled)
classes = np.unique(y_test)
n_classes = len(classes)
y_test_binarized = label_binarize(y_test, classes=classes)

fig, ax = plt.subplots(figsize=(8, 6))

for i in range(n_classes):
    RocCurveDisplay.from_predictions(
        y_test_binarized[:, i],
        y_pred_proba_knn[:, i],
        name=f'ROC curve for class {classes[i]}',
        ax=ax,
    )


ax.plot([0, 1], [0, 1], linestyle='--', lw=2, color='r', label='Chance Level',
ax.set_title('ROC AUC Curve: K-Nearest Neighbors (One-vs-Rest)')
ax.set_xlabel('False Positive Rate')
ax.set_ylabel('True Positive Rate')
ax.grid(linestyle='--', alpha=0.7)
ax.legend(loc='lower right')
plt.show()
```



```
In [ ]: import pickle

output_dir = '/content/drive/MyDrive/AIML_lab_sem4/exp 5,6,7/Exp8_models'

with open(os.path.join(output_dir, 'Logistic_Regression_model.pkl'), 'wb') as file:
    pickle.dump(lr, file)
with open(os.path.join(output_dir, 'Decision_Tree_model.pkl'), 'wb') as file:
    pickle.dump(dt, file)
with open(os.path.join(output_dir, 'Random_Forest_model.pkl'), 'wb') as file:
    pickle.dump(rf, file)
with open(os.path.join(output_dir, 'Gradient_Boosting_model.pkl'), 'wb') as file:
    pickle.dump(gb, file)
with open(os.path.join(output_dir, 'K_Nearest_Neighbors_model.pkl'), 'wb') as file:
    pickle.dump(knn, file)
```



Experiment 9

Physiological signal classification

Aim : Implement a classification model to distinguish between normal and abnormal physiological signals using extracted signal features. (perform comparative analysis on different ML models)

Pipeline :

1.Problem Definition

The primary goal of this experiment is to automate the detection of abnormal heart rhythms (arrhythmias) which is critical for early diagnosis of cardiovascular diseases:

- Classification Task: Categorize heartbeat samples into "Normal" (0) and "Abnormal" (1).
- Objective: Perform a comparative analysis using a Decision Tree, Random Forest, and SVM, and ultimately combine them into a Voting Classifier to improve overall prediction stability and accuracy.

2. Data Understanding

- Dataset Content: The dataset consists of ECG recordings from the PTB Diagnostic Data sets, containing thousands of heartbeat samples.
- Input Features (X): 187 numerical features representing the normalized intensity of the ECG signal over a single heartbeat period.
- Target (y): Binary labels where 0 represents a Normal heartbeat and 1 represents an Abnormal heartbeat

3. Data Preprocessing

To prepare the medical signal data for the machine learning models:

- Data Integration: Combined separate CSV files for normal and abnormal cases into a single unified dataframe.
- Feature Extraction: Separated the signal data (first 187 columns) from the ground truth labels (last column).
- Class Labeling: Explicitly assigned numerical values (0 and 1) to distinguish between the two health states.s.

4. Train-Test Split

To validate the model's diagnostic reliability on unseen patients:

- **Data Allocation:** The combined dataset was split into a Training set (80%) to teach the models and a Test set (20%) for final validation.
- **Random State:** A fixed seed (42) was used to ensure the results are reproducible across different runs.

5. Model Selection and Training

Four distinct configurations were trained to compare individual performance vs. ensemble performance:

- **Decision Tree:** A baseline model using 'entropy' as the split criterion with a maximum depth of 10.
- **Random Forest:** An ensemble of 100 trees to reduce variance and improve accuracy through bagging.
- **Support Vector Machine (SVM):** A high-dimensional classifier used to find the optimal hyperplane between normal and abnormal signals.
- **Voting Classifier (Soft Voting):** A meta-classifier that aggregates the probability predictions of the three models above to make a final "consensus" decision.

6. Predictions

- Each trained model was used to predict the labels for the unseen samples in the test set.
- The predicted data is further used in evaluation of the model.

7. Model Evaluation and Comparison

The performance was evaluated using Precision, Recall, and F1-Score, specifically focusing on the model's ability to identify abnormal cases correctly.

	Accuracy	Precision	Recall	F1
Decision tree	0.8952	0.9355	0.9182	0.9268
Random forest	0.9705	0.9701	0.9895	0.9797
SVM	0.9024	0.9139	0.9548	0.9339
Voting ensemble	0.9443	0.9567	0.9667	0.9617

Auc-roc curve:

Decision Tree	Randon Forest	SVM	Voting Ensemble
0.93	0.99	0.94	0.98

8. Model Saving

The final step ensures the diagnostic tool is ready for clinical application:

- **Serialization:** The trained Voting Classifier was saved into a .pkl file using the Pickle library.
- **Deployment Ready:** This file allows medical software to load the model and analyze live ECG streams without needing to retrain the algorithms.

```
In [1]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
```

```
In [2]: from sklearn.model_selection import train_test_split
from sklearn.metrics import classification_report, confusion_matrix, accuracy_
from sklearn.tree import DecisionTreeClassifier
from sklearn.ensemble import RandomForestClassifier, VotingClassifier, Bagging
from sklearn.svm import SVC
```

```
In [3]: train_df = pd.read_csv('/content/drive/MyDrive/AIML_lab_sem4/Exp 9 , 10/mitbih_
test_df = pd.read_csv('/content/drive/MyDrive/AIML_lab_sem4/Exp 9 , 10/mitbih_
df_normal = pd.read_csv('/content/drive/MyDrive/AIML_lab_sem4/Exp 9 , 10/ptbdb
df_abnormal = pd.read_csv('/content/drive/MyDrive/AIML_lab_sem4/Exp 9 , 10/ptb
```

```
In [4]: df_normal[187] = 0
df_abnormal[187] = 1
```

```
In [5]: df_combined = pd.concat([df_normal, df_abnormal], axis=0)
X = df_combined.iloc[:, :-1].values
y = df_combined.iloc[:, -1].values
```

```
In [6]: X_train, X_test, y_train, y_test = train_test_split(
    X, y, test_size=0.2, random_state=42, stratify=y
)
```

```
In [7]: dt_model = DecisionTreeClassifier(criterion='entropy', max_depth=10, random_st
```

```
In [8]: dt_model.fit(X_train, y_train)
y_pred_dt = dt_model.predict(X_test)
```


```
In [9]: print(classification_report(y_test, y_pred_dt))
```

	precision	recall	f1-score	support
0	0.80	0.84	0.82	809
1	0.94	0.92	0.93	2102
accuracy			0.90	2911
macro avg	0.87	0.88	0.87	2911
weighted avg	0.90	0.90	0.90	2911

```
In [10]: print(f"Accuracy: {accuracy_score(y_test, y_pred_dt):.4f}")
print(f"Precision: {precision_score(y_test, y_pred_dt, zero_division=0):.4f}")
print(f"Recall: {recall_score(y_test, y_pred_dt):.4f}")
print(f"F1-Score: {f1_score(y_test, y_pred_dt):.4f}")
```


Accuracy: 0.8952
Precision: 0.9355
Recall: 0.9182
F1-Score: 0.9268

```
In [11]: cm = confusion_matrix(y_test, y_pred_dt)
plt.figure(figsize=(8, 6))
sns.heatmap(cm, annot=True, fmt='d', cmap='Blues', xticklabels=['Normal', 'Abn
plt.title('Confusion Matrix for Decision Tree Classifier')
plt.xlabel('Predicted Label')
plt.ylabel('True Label')
plt.show()
```



```
In [12]: rf_model = RandomForestClassifier(n_estimators=100, criterion='gini', n_jobs=-
```

```
In [13]: rf_model.fit(X_train, y_train)
```

```
Out[13]: * RandomForestClassifier
RandomForestClassifier(n_jobs=-1, random_state=42)
```

```
In [14]: y_pred_rf = rf_model.predict(X_test)
```


```
In [15]: print(classification_report(y_test, y_pred_rf))
```

	precision	recall	f1-score	support
0	0.97	0.92	0.95	809
1	0.97	0.99	0.98	2102
accuracy			0.97	2911
macro avg	0.97	0.96	0.96	2911
weighted avg	0.97	0.97	0.97	2911

```
In [16]: print(f"Accuracy: {accuracy_score(y_test, y_pred_rf):.4f}")
print(f"Precision: {precision_score(y_test, y_pred_rf, zero_division=0):.4f}")
print(f"Recall: {recall_score(y_test, y_pred_rf):.4f}")
print(f"F1-Score: {f1_score(y_test, y_pred_rf):.4f}")
```

```
Accuracy: 0.9705
Precision: 0.9701
Recall: 0.9895
F1-Score: 0.9797
```

```
In [17]: cm_rf = confusion_matrix(y_test, y_pred_rf)
plt.figure(figsize=(8, 6))
sns.heatmap(cm_rf, annot=True, fmt='d', cmap='Blues', xticklabels=['Normal',
plt.title('Confusion Matrix for Random Forest Classifier')
plt.xlabel('Predicted Label')
plt.ylabel('True Label')
plt.show()
```



```
In [18]: svm_model = SVC(random_state=42)
```

```
In [19]: svm_model.fit(X_train, y_train)
```

```
Out[19]: SVC
SVC(random_state=42)
```

```
In [20]: y_pred_svm = svm_model.predict(X_test)
```


```
In [21]: print(classification_report(y_test, y_pred_svm))
```

	precision	recall	f1-score	support
0	0.87	0.77	0.81	809
1	0.91	0.95	0.93	2102
accuracy			0.90	2911
macro avg	0.89	0.86	0.87	2911
weighted avg	0.90	0.90	0.90	2911

```
In [22]: print(f"Accuracy: {accuracy_score(y_test, y_pred_svm):.4f}")
print(f"Precision: {precision_score(y_test, y_pred_svm, zero_division=0):.4f}")
print(f"Recall: {recall_score(y_test, y_pred_svm):.4f}")
print(f"F1-Score: {f1_score(y_test, y_pred_svm):.4f}")
```


```
Accuracy: 0.9024
Precision: 0.9139
Recall: 0.9548
F1-Score: 0.9339
```

```
In [23]: cm_svm = confusion_matrix(y_test, y_pred_svm)
plt.figure(figsize=(8, 6))
sns.heatmap(cm_svm, annot=True, fmt='d', cmap='Blues', xticklabels=['Normal',
plt.title('Confusion Matrix for Support Vector Machine Classifier')
plt.xlabel('Predicted Label')
plt.ylabel('True Label')
plt.show()
```



```
In [24]: voting_classifier = VotingClassifier(estimators=[
    ('dt', DecisionTreeClassifier(criterion='entropy', max_depth=10, random_st
    ('rf', RandomForestClassifier(n_estimators=100, criterion='gini', n_jobs=-
    ('svm', SVC(random_state=42, probability=True))
], voting='soft', n_jobs=-1)
```

```
In [25]: voting_classifier.fit(X_train, y_train)
```



```
In [26]: y_pred_voting = voting_classifier.predict(X_test)
```


```
In [27]: print(classification_report(y_test, y_pred_voting))
```

	precision	recall	f1-score	support
0	0.91	0.89	0.90	809
1	0.96	0.97	0.96	2102
accuracy			0.94	2911
macro avg	0.93	0.93	0.93	2911
weighted avg	0.94	0.94	0.94	2911

```
In [28]: print(f"Accuracy: {accuracy_score(y_test, y_pred_voting):.4f}")
print(f"Precision: {precision_score(y_test, y_pred_voting, zero_division=0):.4f}")
print(f"Recall: {recall_score(y_test, y_pred_voting):.4f}")
print(f"F1-Score: {f1_score(y_test, y_pred_voting):.4f}")
```

```
Accuracy: 0.9443
Precision: 0.9567
Recall: 0.9667
F1-Score: 0.9617
```


```
In [29]: cm_voting = confusion_matrix(y_test, y_pred_voting)
plt.figure(figsize=(8, 6))
sns.heatmap(cm_voting, annot=True, fmt='d', cmap='Blues', xticklabels=['Normal', 'Abnormal'])
plt.title('Confusion Matrix for Voting Classifier')
plt.xlabel('Predicted Label')
plt.ylabel('True Label')
plt.show()
```



```
In [36]: plt.figure(figsize=(10, 8))

plt.plot(fpr_dt, tpr_dt, color='darkorange', lw=2, label=f'Decision Tree (AUC = {roc_auc_dt:.2f})')
plt.plot(fpr_rf, tpr_rf, color='green', lw=2, label=f'Random Forest (AUC = {roc_auc_rf:.2f})')
plt.plot(fpr_svm, tpr_svm, color='red', lw=2, label=f'SVM (AUC = {roc_auc_svm:.2f})')
plt.plot(fpr_voting, tpr_voting, color='purple', lw=2, label=f'Voting Classifier (AUC = {roc_auc_voting:.2f})')


plt.plot([0, 1], [0, 1], color='navy', lw=2, linestyle='--', label='Random Classifier')
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('ROC Curve for All Classifiers')
plt.legend(loc='lower right')
plt.grid(True)
plt.show()
```

```
In [30]: import pickle

filename = '/content/drive/MyDrive/AIML_lab_sem4/Exp 9 , 10/voting_classifier_

with open(filename, 'wb') as file:
    pickle.dump(voting_classifier, file)
```



Experiment - 10

Iris Classification

Aim : Iris Flower Classification: Use the Iris dataset to build a classification model that predicts the species of iris flowers. [Dataset: Load dataset from sk-learn] .

Pipeline :

1. Problem Definition

The primary goal of this experiment is to develop a robust automated classification system to identify biological species based on physical measurements:

- **Classification Task:** We must determine if a data point belongs to the *Setosa*, *Versicolor*, or *Virginica* class.
- **Objective:** We perform a detailed comparative analysis across three distinct algorithms Decision Tree, K-Nearest Neighbors (KNN), and Logistic Regression to identify the most accurate architecture for multiclass classification.

2. Data Understanding

- **Dataset Content:** We utilize the classic Iris dataset consisting of 150 samples with four physical attributes: sepal length, sepal width, petal length, and petal width.
- **Input Features (X):** We analyze numerical data representing the four flower measurements.
- **Target (y):** We use labels where 0, 1, and 2 represent the three specific species of iris.

3. Data Preprocessing

To ensure the data is suitable for algorithmic processing, we perform the following:

- **Data Formatting:** We separate the dataset into a feature matrix (x) and a target vector (y).
- **Cleanliness:** As the Scikit-Learn iris dataset is a standard benchmark, the data is verified to be clean and correctly formatted for supervised learning.

4. Train-Test Split

To evaluate how well the system identifies species in new samples, we strategically divide the dataset:

- **Data Allocation:** We split the dataset into a Training set (80%) to allow the models to learn the patterns and a Test set (20%) to validate accuracy.
- **Reproducibility:** We use a random_state of 42 to ensure that the data split remains consistent across different experimental runs.

5. Model Selection and Training

We initialize and train three different classification models on the same data to ensure a fair comparison:

- **Decision Tree Classifier:** We use this tree-based model that breaks down the dataset into smaller subsets while at the same time an associated decision tree is incrementally developed.
- **K-Nearest Neighbors (KNN):** We implement this instance-based method that classifies a point based on the majority class of its nearest neighbors in the feature space.
- **Logistic Regression:** We task this linear model with learning the probability of a sample belonging to each of the three species classes.

6. Predictions

- **Execution:** Once training is complete, we task each model with classifying the 30 unseen samples in our test set.
- **Output Generation:** We generate prediction arrays where each algorithm flags points as either class 0, 1, or 2.

7. Model Evaluation and Comparison

The performance of each model is measured using a Confusion Matrix for visual verification and a Classification Report containing Precision, Recall, and F1-Score. Based on the experimental outcomes:

	Accuracy	Precision	Recall	F1
Decision Tree	1.00	1.00	1.00	1.00
KNN	1.00	1.00	1.00	1.00
Logistic Regression	1.00	1.00	1.00	1.00

As the data is very simple and got perfect scores in all the models the auc-roc of all the models was also 1.00.

8. Model Saving

The final step is ensuring our models are persistent for future deployment:

- **Implementation:** We save the top-performing trained model (Logistic Regression) using the Pickle library to serialize the object into a permanent .pkl file.
- **Final Output:** We generate the file `logistic_iris.pkl` to allow for immediate deployment in an automated classification system without the need for retraining.

```
In [1]: from sklearn.datasets import load_iris
        from sklearn.model_selection import train_test_split
        from sklearn.tree import DecisionTreeClassifier
        from sklearn.metrics import accuracy_score, classification_report
```

```
In [2]: iris = load_iris()
        X = iris.data
        y = iris.target
```

```
In [3]: X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)
```

```
In [4]: dt_model = DecisionTreeClassifier()
```


```
In [5]: dt_model.fit(X_train, y_train)
```

```
Out[5]: +DecisionTreeClassifier
        DecisionTreeClassifier()
```

```
In [6]: y_pred = dt_model.predict(X_test)
```

```
In [7]: import matplotlib.pyplot as plt
        import seaborn as sns
        from sklearn.metrics import confusion_matrix

        cm = confusion_matrix(y_test, y_pred)
        plt.figure(figsize=(8, 6))
        sns.heatmap(cm, annot=True, fmt='d', cmap='Blues',
                    xticklabels=iris.target_names,
                    yticklabels=iris.target_names)
        plt.xlabel('Predicted Labels')
        plt.ylabel('Actual Labels')
        plt.title('Confusion Matrix: Decision Tree Classifier')
        plt.show()
```



```
In [8]: from sklearn.metrics import classification_report
```

```
In [9]: print(classification_report(y_test, y_pred, target_names=iris.target_names))
```

	precision	recall	f1-score	support
setosa	1.00	1.00	1.00	10
versicolor	1.00	1.00	1.00	9
virginica	1.00	1.00	1.00	11
accuracy			1.00	30
macro avg	1.00	1.00	1.00	30
weighted avg	1.00	1.00	1.00	30

```
In [10]: from sklearn.neighbors import KNeighborsClassifier
```


```
knn_model = KNeighborsClassifier()
knn_model.fit(X_train, y_train)
```

```
Out[10]: KNeighborsClassifier
```

```
KNeighborsClassifier()
```

```
In [11]: y_pred_knn = knn_model.predict(X_test)
```

```
In [12]: cm_knn = confusion_matrix(y_test, y_pred_knn)
plt.figure(figsize=(8, 6))
sns.heatmap(cm_knn, annot=True, fmt='d', cmap='Greens',
            xticklabels=iris.target_names,
            yticklabels=iris.target_names)
plt.xlabel('Predicted Labels')
plt.ylabel('Actual Labels')
plt.title('Confusion Matrix: K-Nearest Neighbors Classifier')
plt.show()
```



```
In [13]: print(classification_report(y_test, y_pred_knn, target_names=iris.target_names))
```


	precision	recall	f1-score	support
setosa	1.00	1.00	1.00	10
versicolor	1.00	1.00	1.00	9
virginica	1.00	1.00	1.00	11
accuracy			1.00	30
macro avg	1.00	1.00	1.00	30
weighted avg	1.00	1.00	1.00	30

```
In [14]: from sklearn.linear_model import LogisticRegression
logreg_model = LogisticRegression(max_iter=200, random_state=42)
logreg_model.fit(X_train, y_train)
```

```
Out[14]: LogisticRegression
LogisticRegression(max_iter=200, random_state=42)
```

```
In [15]: y_pred_logreg = logreg_model.predict(X_test)
```

```
In [16]: cm_logreg = confusion_matrix(y_test, y_pred_logreg)
plt.figure(figsize=(8, 6))
sns.heatmap(cm_logreg, annot=True, fmt='d', cmap='Purples',
            xticklabels=iris.target_names,
            yticklabels=iris.target_names)
plt.xlabel('Predicted Labels')
plt.ylabel('Actual Labels')
plt.title('Confusion Matrix: Logistic Regression Classifier')
plt.show()
```



```
In [17]: print(classification_report(y_test, y_pred_logreg, target_names=iris.target_names))
```

	precision	recall	f1-score	support
setosa	1.00	1.00	1.00	10
versicolor	1.00	1.00	1.00	9
virginica	1.00	1.00	1.00	11
accuracy			1.00	30
macro avg	1.00	1.00	1.00	30
weighted avg	1.00	1.00	1.00	30

```
In [18]: from sklearn.metrics import roc_curve, auc
from sklearn.preprocessing import label_binarize
import numpy as np


# Binarize the output for ROC AUC calculation (one-vs-rest)
y_test_bin = label_binarize(y_test, classes=np.unique(y))
n_classes = y_test_bin.shape[1]
```

ROC AUC Curve for Decision Tree Classifier

```
In [20]: y_prob_dt = dt_model.predict_proba(X_test)

plt.figure(figsize=(10, 8))
colors = ['red', 'green', 'blue']
for i in range(n_classes):
    fpr, tpr, _ = roc_curve(y_test_bin[:, i], y_prob_dt[:, i])
    roc_auc = auc(fpr, tpr)
    plt.plot(fpr, tpr, color=colors[i], lw=2, label=f'ROC curve of class {iris

plt.plot([0, 1], [0, 1], 'k--', lw=2)
plt.xlim([0.0, 1.0])
plt.ylim([0.0, 1.05])
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('Receiver Operating Characteristic (ROC) - Decision Tree Classifier')
plt.legend(loc="lower right")
plt.grid(True)
plt.show()
```




ROC AUC Curve for K-Nearest Neighbors Classifier

```
In [21]: y_prob_knn = knn_model.predict_proba(X_test)

plt.figure(figsize=(10, 8))
colors = ['red', 'green', 'blue']
for i in range(n_classes):
    fpr, tpr, _ = roc_curve(y_test_bin[:, i], y_prob_knn[:, i])
    roc_auc = auc(fpr, tpr)
    plt.plot(fpr, tpr, color=colors[i], lw=2, label=f'ROC curve of class {iris

plt.plot([0, 1], [0, 1], 'k--', lw=2)
plt.xlim([0.0, 1.0])
plt.ylim([0.0, 1.05])
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('Receiver Operating Characteristic (ROC) - K-Nearest Neighbors Class
plt.legend(loc="lower right")
plt.grid(True)
plt.show()
```




ROC AUC Curve for Logistic Regression Classifier

```
In [22]: y_prob_logreg = logreg_model.predict_proba(X_test)


plt.figure(figsize=(10, 8))
colors = ['red', 'green', 'blue']
for i in range(n_classes):
    fpr, tpr, _ = roc_curve(y_test_bin[:, i], y_prob_logreg[:, i])
    roc_auc = auc(fpr, tpr)
    plt.plot(fpr, tpr, color=colors[i], lw=2, label=f'ROC curve of class {iris

plt.plot([0, 1], [0, 1], 'k--', lw=2)
plt.xlim([0.0, 1.0])
plt.ylim([0.0, 1.05])
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('Receiver Operating Characteristic (ROC) - Logistic Regression Class
plt.legend(loc="lower right")
plt.grid(True)
plt.show()
```



```
In [18]: import pickle
```

```
model_filename = '/content/drive/MyDrive/AIML_lab_sem4/Exp 9 , 10/logistic_iri  
with open(model_filename, 'wb') as file:  
    pickle.dump(logreg_model, file)
```



Experiment 11

Diabetes prediction using bagging classifier

Aim = To predict whether a person has diabetes based on features such as blood pressure, skin thickness, age, etc., using the bagging ensemble technique. Also perform comparative analysis among the bagging classifier, random forest, and the decision tree classifier.

Pipeline =>

1. Problem Definition

The goal of this experiment is to implement an automated classification system to predict student diabetes outcomes based on diagnostic metrics.

- **Classification Task:** This is a binary classification problem where samples are categorized into two distinct classes (Outcome 0 or 1).
- **Objective:** Perform a comparative analysis between a single Decision Tree and ensemble methods, specifically Random Forest and Bagging, to determine which architecture offers the most reliable predictive performance.

2. Data Understanding

The experiment utilizes a diabetes dataset containing several medical indicator features:

- **Input Features (X):** Pregnancies, Glucose, Blood Pressure, Skin Thickness, Insulin, BMI, Diabetes Pedigree Function, and Age.
- **Target (y):** Outcome (Binary labels where 1 indicates the presence of diabetes and 0 indicates absence).

3. Data Preprocessing

To prepare the data for training, the following cleaning and transformation steps were executed:

- **Missing Value Handling:** Columns containing invalid zero values (Glucose, BloodPressure, SkinThickness, Insulin, BMI) were identified. These zeros were replaced with NaN and subsequently imputed using the median value of each column.
- **Feature-Target Separation:** The dataset was divided into input features (X) and the target variable (y).

- **Standardization:** The StandardScaler was applied to normalize the feature set, ensuring that variables with different scales do not bias the algorithms.

4. Train-Test Split

- **Data Allocation:** The dataset was partitioned into a Training set (80%) to build the models and a Test set (20%) to validate performance on unseen data.
- **Random State:** A fixed random_state=42 was used to ensure reproducibility of the data split and model results.

5. Model Selection and Training

Three distinct classification architectures were initialized and trained:

- **Decision Tree:** A base non-linear model that creates decision rules based on feature thresholds.
- **Random Forest:** An ensemble technique that builds multiple decision trees on different data subsets and averages their results to improve accuracy and control overfitting.
- **Bagging Classifier:** An ensemble method that uses a DecisionTreeClassifier as the base estimator, training 200 different versions of the model on boosted samples to reduce variance.
- The all three models are trained on the training data.

6. Predictions

- The standardized test set was processed by all three trained models to generate predicted labels for the student samples.
- The models transformed medical input patterns into binary classifications (0 or 1), which were then compared against the actual ground-truth outcomes from the test set.

7. Model Evaluation and Comparison

The models were evaluated using a suite of classification metrics and visualization tools:

	Accuracy	Precision	Recall	F1
Decision Tree	0.71	0.59	0.61	0.60
Random Forest	0.74	0.63	0.65	0.64
Bagging(Decision Tree)	0.77	0.67	0.70	0.69

Auc-Roc Score:

Decision Tree	Random Forest	Bagging
0.69	0.83	0.83

By looking at the results we can conclude that the Bagging classifier out performs random forest and decision tree in all the metrics.

8. Model Saving

To ensure the models are ready for deployment or future use without retraining:

- **Serialization:** All three trained models (Decision Tree, Random Forest, and Bagging Classifier) were exported into .pkl files using the pickle library.
- **Utility:** These files allow for the immediate reloading of the predictive system into a production environment for real-time analysis.

```
In [2]: import pandas as pd
        from sklearn.preprocessing import StandardScaler
        from sklearn.metrics import classification_report
        import numpy as np
        import matplotlib.pyplot as plt
        from sklearn.model_selection import train_test_split
        from sklearn.tree import DecisionTreeClassifier
        from sklearn.ensemble import BaggingClassifier, RandomForestClassifier
        from sklearn.metrics import accuracy_score, precision_score, recall_score, f1_
```

```
In [3]: df = pd.read_csv('/content/drive/MyDrive/AIML_lab_sem4/exp 11,12/dataset_diabe
```

```
In [4]: df.head()
```

```
Out[4]:
```

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI	DiabetesF
0	6	148	72	35	0	33.6	
1	1	85	66	29	0	26.6	
2	8	183	64	0	0	23.3	
3	1	89	66	23	94	28.1	
4	0	137	40	35	168	43.1	

```
In [5]: df.shape
```

```
Out[5]: (768, 9)
```

```
In [6]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 768 entries, 0 to 767
Data columns (total 9 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   Pregnancies                          768 non-null    int64
1   Glucose                              768 non-null    int64
2   BloodPressure                        768 non-null    int64
3   SkinThickness                       768 non-null    int64
4   Insulin                             768 non-null    int64
5   BMI                                 768 non-null    float64
6   DiabetesPedigreeFunction             768 non-null    float64
7   Age                                 768 non-null    int64
8   Outcome                             768 non-null    int64
dtypes: float64(2), int64(7)
memory usage: 54.1 KB
```

```
In [7]: cols_with_zeros = ['Glucose', 'BloodPressure', 'SkinThickness', 'Insulin', 'BMI']
        df[cols_with_zeros] = df[cols_with_zeros].replace(0, np.nan)
        df.fillna(df.median(), inplace=True)
```

```
In [8]: scaler = StandardScaler()
X = scaler.fit_transform(df.drop('Outcome', axis=1))
y = df['Outcome']
```

```
In [9]: X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)
```

```
In [10]: dt = DecisionTreeClassifier(random_state=42)
```

```
In [11]: dt.fit(X_train, y_train)
```

```
Out[11]: DecisionTreeClassifier
DecisionTreeClassifier(random_state=42)
```

```
In [12]: y_pred_dt = dt.predict(X_test)
```

```
In [13]: accuracy = accuracy_score(y_test, y_pred_dt)
precision = precision_score(y_test, y_pred_dt)
recall = recall_score(y_test, y_pred_dt)
f1 = f1_score(y_test, y_pred_dt)

print(f"Accuracy: {accuracy:.4f}")
print(f"Precision: {precision:.4f}")
print(f"Recall: {recall:.4f}")
print(f"F1 Score: {f1:.4f}")
```

```
Accuracy: 0.7143
Precision: 0.5965
Recall: 0.6182
F1 Score: 0.6071
```

```
In [14]: print(classification_report(y_test, y_pred_dt))
```

	precision	recall	f1-score	support
0	0.78	0.77	0.78	99
1	0.60	0.62	0.61	55
accuracy			0.71	154
macro avg	0.69	0.69	0.69	154
weighted avg	0.72	0.71	0.72	154

```
In [15]: from sklearn.metrics import confusion_matrix
cm_dt = confusion_matrix(y_test, y_pred_dt)
```

```
In [16]: import seaborn as sns
import matplotlib.pyplot as plt
```


```
In [17]: plt.figure(figsize=(8, 6))
sns.heatmap(cm_dt, annot=True, fmt='d', cmap='Blues', cbar=False,
```



```

xticklabels=['Predicted 0', 'Predicted 1'],
yticklabels=['Actual 0', 'Actual 1'])
plt.xlabel('Predicted Label')
plt.ylabel('True Label')
plt.title('Confusion Matrix for Decision Tree Classifier')
plt.show()

```



```

In [18]: rf = RandomForestClassifier(n_estimators = 100 ,random_state=42)
         rf.fit(X_train, y_train)

```

```

Out[18]: RandomForestClassifier
RandomForestClassifier(random_state=42)

```

```

In [19]: y_pred_rf = rf.predict(X_test)

```

```

In [20]: cm_rf = confusion_matrix(y_test, y_pred_rf)

```

```

In [21]: accuracy_rf = accuracy_score(y_test, y_pred_rf)
         precision_rf = precision_score(y_test, y_pred_rf)
         recall_rf = recall_score(y_test, y_pred_rf)

```

```

f1_rf = f1_score(y_test, y_pred_rf)

print("Random Forest Classifier Metrics:")
print(f"Accuracy: {accuracy_rf:.4f}")
print(f"Precision: {precision_rf:.4f}")
print(f"Recall: {recall_rf:.4f}")
print(f"F1 Score: {f1_rf:.4f}")

```

```

Random Forest Classifier Metrics:
Accuracy: 0.7403
Precision: 0.6316
Recall: 0.6545
F1 Score: 0.6429

```

```

In [22]: print("\nClassification Report for Random Forest Classifier:")
         print(classification_report(y_test, y_pred_rf))

```

```

Classification Report for Random Forest Classifier:


```

	precision	recall	f1-score	support
0	0.80	0.79	0.80	99
1	0.63	0.65	0.64	55
accuracy			0.74	154
macro avg	0.72	0.72	0.72	154
weighted avg	0.74	0.74	0.74	154

```

In [23]: plt.figure(figsize=(8, 6))
         sns.heatmap(cm_rf, annot=True, fmt='d', cmap='Greens', cbar=False,
                    xticklabels=['Predicted 0', 'Predicted 1'],
                    yticklabels=['Actual 0', 'Actual 1'])
         plt.xlabel('Predicted Label')
         plt.ylabel('True Label')
         plt.title('Confusion Matrix for Random Forest Classifier')
         plt.show()

```



```
In [24]: bagging_classifier = BaggingClassifier(estimator=DecisionTreeClassifier(random  
bagging_classifier.fit(X_train, y_train)
```

```
Out[24]:
```

- BaggingClassifier
 - estimator:
DecisionTreeClassifier
 - DecisionTreeClassifier

```
In [25]: y_pred_bagging = bagging_classifier.predict(X_test)
```

```
In [26]: cm_bagging = confusion_matrix(y_test, y_pred_bagging)
```

```
In [27]: accuracy_bagging = accuracy_score(y_test, y_pred_bagging)  
precision_bagging = precision_score(y_test, y_pred_bagging)  
recall_bagging = recall_score(y_test, y_pred_bagging)  
f1_bagging = f1_score(y_test, y_pred_bagging)
```

```

print("Bagging Classifier Metrics (50 Decision Trees):")
print(f"Accuracy: {accuracy_bagging:.4f}")
print(f"Precision: {precision_bagging:.4f}")
print(f"Recall: {recall_bagging:.4f}")
print(f"F1 Score: {f1_bagging:.4f}")

```

```

Bagging Classifier Metrics (50 Decision Trees):
Accuracy: 0.7727
Precision: 0.6724
Recall: 0.7091
F1 Score: 0.6903

```

```

In [28]: print("\nClassification Report for Bagging Classifier:")
print(classification_report(y_test, y_pred_bagging))

```

```

Classification Report for Bagging Classifier:


```

	precision	recall	f1-score	support
0	0.83	0.81	0.82	99
1	0.67	0.71	0.69	55
accuracy			0.77	154
macro avg	0.75	0.76	0.76	154
weighted avg	0.78	0.77	0.77	154

```

In [29]: plt.figure(figsize=(8, 6))
sns.heatmap(cn_bagging, annot=True, fmt='d', cmap='Purples', cbar=False,
            xticklabels=['Predicted 0', 'Predicted 1'],
            yticklabels=['Actual 0', 'Actual 1'])
plt.xlabel('Predicted Label')
plt.ylabel('True Label')
plt.title('Confusion Matrix for Bagging Classifier')
plt.show()

```



```
In [30]: from sklearn.metrics import roc_curve, auc, RocCurveDisplay

plt.figure(figsize=(10, 8))


probs_dt = dt.predict_proba(X_test)[:, 1]
fpr_dt, tpr_dt, _ = roc_curve(y_test, probs_dt)
roc_auc_dt = auc(fpr_dt, tpr_dt)
RocCurveDisplay(fpr=fpr_dt, tpr=tpr_dt, roc_auc=roc_auc_dt, estimator_name='Decision Tree')

probs_rf = rf.predict_proba(X_test)[:, 1]
fpr_rf, tpr_rf, _ = roc_curve(y_test, probs_rf)
roc_auc_rf = auc(fpr_rf, tpr_rf)
RocCurveDisplay(fpr=fpr_rf, tpr=tpr_rf, roc_auc=roc_auc_rf, estimator_name='Random Forest')

probs_bagging = bagging_classifier.predict_proba(X_test)[:, 1]
fpr_bagging, tpr_bagging, _ = roc_curve(y_test, probs_bagging)
roc_auc_bagging = auc(fpr_bagging, tpr_bagging)
RocCurveDisplay(fpr=fpr_bagging, tpr=tpr_bagging, roc_auc=roc_auc_bagging, estimator_name='Bagging Classifier')

plt.plot([0, 1], [0, 1], 'k--', label='Random Chance')
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('ROC Curve Comparison')
```

```
plt.legend(loc='lower right')
plt.grid(True)
plt.show()
```




```
In [34]: selected_models = ['Decision Tree', 'Random Forest', 'Bagging (Decision Tree)']
filtered_metrics_df = metrics_df[metrics_df['Model'].isin(selected_models)]

plt.figure(figsize=(12, 6))
sns.barplot(x='Model', y='Score', hue='Metric', data=filtered_metrics_df, palette='magma')
plt.title('Comparison of Accuracy, F1 Score, Precision, and Recall for Selected Models')
plt.ylabel('Score')
plt.ylim(0.5, 0.85)
plt.legend(title='Metric', bbox_to_anchor=(1.05, 1), loc='upper left')
plt.grid(axis='y', linestyle='--', alpha=0.7)
plt.xticks(rotation=45, ha='right')
```



```
plt.tight_layout()
plt.show()
```



```
In [ ]: import pickle
with open('/content/drive/MyDrive/AIML_lab_sem4/exp 11,12/Exp11_models/decision_tree.pkl', 'wb') as file:
    pickle.dump(dt, file)
```

```
In [ ]: with open('/content/drive/MyDrive/AIML_lab_sem4/exp 11,12/Exp11_models/random_forest.pkl', 'wb') as file:
    pickle.dump(rf, file)
```

```
In [ ]: with open('/content/drive/MyDrive/AIML_lab_sem4/exp 11,12/Exp11_models/bagging_decision_tree.pkl', 'wb') as file:
    pickle.dump(bagging_classifier, file)
```

Experiment 12

Regularization

Aim : Implement L1 (Lasso) and L2 (Ridge) regularization on the given Melbourne House Price Dataset.

Pipeline

1. Problem Definition

The primary goal is to develop a predictive model for real estate valuation using advanced regression techniques:

- **Regression Task:** We aim to predict the continuous numerical value of house prices based on various structural and locational features.
- **Objective:** We perform a comparative analysis between Standard Linear Regression, Lasso (L1), and Ridge (L2) regularization to determine which approach best handles multicollinearity and prevents model over-complexity.

2. Data Understanding

- **Dataset Content:** We utilize the Melbourne House Price dataset, which contains records of real estate transactions.
- **Input Features (X):**
 - Numerical: Rooms, Distance, Bedroom2, Bathroom, Car, Landsize, BuildingArea, YearBuilt, Propertycount.
 - Categorical: Suburb, Address, Type, Method, SellerG, Date, CouncilArea, Regionname.
- **Target (y):** Price (The market value of the property in AUD)

3. Data Preprocessing

To ensure the dataset is optimized for regularization, a rigorous transformation process is applied:

- **Handling Missing Values:** We address null entries in features like 'Car', 'BuildingArea', and 'YearBuilt' using zero-filling or mean/median imputation.
- **Categorical Encoding:** We apply **One-Hot Encoding** (via `get_dummies`) to transform non-numeric features into binary columns compatible with regression algorithms.
- **Feature-Target Separation:** The dataset is divided into the feature matrix (X) and the target variable (y).

4. Train-Test Split

- **Data Allocation:** The dataset is partitioned into a Training set (80%) to fit the model and a Test set (20%) to evaluate its performance on unseen data.
- **Random State:** A fixed random_state is used to ensure the results and data splits are reproducible across different runs.

5. Model Selection and Training

We initialize and train three different regression architectures to observe the impact of penalties:

- **Linear Regression:** A baseline model that calculates coefficients without any regularization constraints.
- **Lasso Regression (L1):** A model that adds an absolute value penalty to the loss function, which can shrink some coefficients to exactly zero, performing automatic feature selection.
- **Ridge Regression (L2):** A model that adds a squared magnitude penalty to the loss function, which helps manage multicollinearity by shrinking coefficients but not setting them to zero.

6. Predictions

- **Execution:** The trained models (Linear, Lasso, and Ridge) are used to predict house prices for the properties in the test set.
- **Storage:** Predicted values are stored to calculate error metrics and compare against the actual house price

7. Model Evaluation and Comparison

The models are evaluated using the R2 score (Coefficient of Determination) to measure how well the models generalize

Linear Regression	Ridge	Lasso
0.6426	0.6529	0.6530

8. Model Saving

- **Implementation:** The optimized regularized model and any preprocessing objects are serialized into .pkl files using the **Pickle** library.
- **Final Output:** These files allow the prediction system to be reloaded for real-time house price estimations without needing to retrain the model

```
In [17]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.model_selection import train_test_split
from sklearn.impute import SimpleImputer, KNNImputer
from sklearn.preprocessing import StandardScaler, OneHotEncoder, OrdinalEncoder
from sklearn.linear_model import LinearRegression, Lasso, Ridge
from sklearn.metrics import r2_score

df = pd.read_csv('/content/drive/MyDrive/AIML_lab_sem4/exp 11,12/dataset_Melbo
```

```
In [18]: df.head()
```

```
Out[18]:
```

	Suburb	Address	Rooms	Type	Price	Method	SellerG	Date	Distar
0	Abbotsford	68 Studley St	2	h	NaN	SS	Jellis	3/09/2016	
1	Abbotsford	85 Turner St	2	h	1480000.0	S	Biggin	3/12/2016	
2	Abbotsford	25 Bloomburg St	2	h	1035000.0	S	Biggin	4/02/2016	
3	Abbotsford	18/659 Victoria St	3	u	NaN	VB	Rounds	4/02/2016	
4	Abbotsford	5 Charles St	3	h	1465000.0	SP	Biggin	4/03/2017	

5 rows x 21 columns

```
In [19]: df.info()
```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 34857 entries, 0 to 34856
Data columns (total 21 columns):
#   Column              Non-Null Count  Dtype
---  -
0   Suburb              34857 non-null  object
1   Address             34857 non-null  object
2   Rooms              34857 non-null  int64
3   Type               34857 non-null  object
4   Price              27247 non-null  float64
5   Method             34857 non-null  object
6   SellerG            34857 non-null  object
7   Date               34857 non-null  object
8   Distance           34856 non-null  float64
9   Postcode           34856 non-null  float64
10  Bedroom2           26640 non-null  float64
11  Bathroom            26631 non-null  float64
12  Car                 26129 non-null  float64
13  Landsize            23047 non-null  float64
14  BuildingArea        13742 non-null  float64
15  YearBuilt           15551 non-null  float64
16  CouncilArea         34854 non-null  object
17  Lattitude           26881 non-null  float64
18  Longitude           26881 non-null  float64
19  Regionname          34854 non-null  object
20  Propertycount       34854 non-null  float64
dtypes: float64(12), int64(1), object(8)
memory usage: 5.6+ MB

```

```
In [20]: df.isnull().sum()
```

```
Out[20]:
```

	0
Suburb	0
Address	0
Rooms	0
Type	0
Price	7610
Method	0
SellerG	0
Date	0
Distance	1
Postcode	1
Bedroom2	8217
Bathroom	8226
Car	8728
Landsize	11810
BuildingArea	21115
YearBuilt	19306
CouncilArea	3
Lattitude	7976
Longtitude	7976
Regionname	3
Propertycount	3

dtype: int64

```
In [21]: df = df.dropna(subset=['Price'])
```

```
In [22]: num_cols = ['Rooms', 'Distance', 'Postcode', 'Bedroom2', 'Bathroom', 'Car', 'Landsize', 'BuildingArea', 'YearBuilt']
cat_cols = ['Type', 'Method', 'CouncilArea', 'Regionname']
```

```
In [23]: X = df[num_cols + cat_cols]
y = df['Price']
```

```
In [24]: df[num_cols] = df[num_cols].fillna(df[num_cols].median())
```

```
In [25]: for col in cat_cols:
df[col] = df[col].fillna(df[col].mode()[0])
```

```
In [26]: df = df.drop(['Address', 'Date'], axis=1)
```

```
df = pd.get_dummies(df, drop_first=True)
```

```
In [27]: X = df.drop('Price', axis=1)
y = df['Price']
```

```
In [28]: X_train, X_test, y_train, y_test = train_test_split(
    X, y, test_size=0.2, random_state=42
)
```

```
In [29]: from sklearn.preprocessing import StandardScaler

scaler = StandardScaler()

X_train_scaled = scaler.fit_transform(X_train)
X_test_scaled = scaler.transform(X_test)
```

```
In [29]:
```

```
In [33]: from sklearn.metrics import mean_squared_error

# Apply Ridge (L2 Regularization)
ridge = Ridge(alpha=100)
ridge.fit(X_train_scaled, y_train)

y_pred_ridge = ridge.predict(X_test_scaled)

print("\nRidge Results:")
mse_ridge = mean_squared_error(y_test, y_pred_ridge)
rmse_ridge = np.sqrt(mse_ridge)
print("MSE:", mse_ridge)
print("RMSE:", rmse_ridge)
print("R2 Score:", r2_score(y_test, y_pred_ridge))
```

```
Ridge Results:
MSE: 148308986100.42667
RMSE: 385109.05741156836
R2 Score: 0.6529238980216601
```

```
In [34]: from sklearn.linear_model import LinearRegression
from sklearn.metrics import mean_squared_error, r2_score

# Apply Linear Regression
linear_reg = LinearRegression()
linear_reg.fit(X_train_scaled, y_train)

y_pred_linear = linear_reg.predict(X_test_scaled)

print("\nLinear Regression Results:")
mse_linear = mean_squared_error(y_test, y_pred_linear)
rmse_linear = np.sqrt(mse_linear)
```

```
print("MSE:", mse_linear)
print("RMSE:", rmse_linear)
print("R2 Score:", r2_score(y_test, y_pred_linear))
```

Linear Regression Results:
MSE: 152722264764.61996
RMSE: 390796.9610483428
R2 Score: 0.6425958417387105

```
In [37]: from sklearn.linear_model import Lasso

# Apply Lasso (L1 Regularization)
lasso = Lasso(alpha=5.0) # You can tune alpha as needed
lasso.fit(X_train_scaled, y_train)

y_pred_lasso = lasso.predict(X_test_scaled)

print("\nLasso Results:")
mse_lasso = mean_squared_error(y_test, y_pred_lasso)
rmse_lasso = np.sqrt(mse_lasso)
print("MSE:", mse_lasso)
print("RMSE:", rmse_lasso)
print("R2 Score:", r2_score(y_test, y_pred_lasso))
```

Lasso Results:
MSE: 148275760119.93
RMSE: 385065.91659082216
R2 Score: 0.6530016542257742

```
/usr/local/lib/python3.12/dist-packages/sklearn/linear_model/_coordinate_descent.py:695: ConvergenceWarning: Objective did not converge. You might want to increase the number of iterations, check the scale of the features or consider increasing regularisation. Duality gap: 5.717e+14, tolerance: 8.882e+11
  model = cd_fast.enet_coordinate_descent(
```

```
In [36]: lasso_coeff = pd.Series(lasso.coef_, index=X.columns)
ridge_coeff = pd.Series(ridge.coef_, index=X.columns)

print("\nNumber of features used by Lasso:", sum(lasso_coeff != 0))
print("Number of features used by Ridge:", len(ridge_coeff))
```

Number of features used by Lasso: 721
Number of features used by Ridge: 749

```
In [41]: r2_linear = r2_score(y_test, y_pred_linear)
r2_ridge = r2_score(y_test, y_pred_ridge)
r2_lasso = r2_score(y_test, y_pred_lasso)

models = ['Linear Regression', 'Ridge', 'Lasso']
r2_scores = [r2_linear, r2_ridge, r2_lasso]

plt.figure(figsize=(10, 6))
ax = sns.barplot(x=models, y=r2_scores, palette='viridis')
plt.title('Comparison of R2 Scores for Different Models')
plt.xlabel('Model')
```



```

plt.ylabel('R2 Score')
plt.ylim(min(r2_scores) * 0.95, max(r2_scores) * 1.05)
plt.grid(axis='y', linestyle='--', alpha=0.7)
for p in ax.patches:
    ax.annotate(f'{p.get_height():.4f}',
                (p.get_x() + p.get_width() / 2., p.get_height()),
                ha='center', va='center',
                xytext=(0, 9),
                textcoords='offset points')


plt.show()

```

/tmp/ipykernel_17265/2865464898.py:9: FutureWarning:

Passing 'palette' without assigning 'hue' is deprecated and will be removed in v0.14.0. Assign the 'x' variable to 'hue' and set 'legend=False' for the same effect.


```
ax = sns.barplot(x=models, y=r2_scores, palette='viridis')
```




```

In [43]: plt.figure(figsize=(12, 6))
plt.plot(y_test.values[:50], label='Actual Values', alpha=0.7)
plt.plot(y_pred_linear[:50], label='Predicted Values (Linear Regression)', alpha=0.7)
plt.title('Actual vs. Predicted Values (Linear Regression) - First 50 Samples')
plt.xlabel('Sample Index')
plt.ylabel('Price')
plt.legend()
plt.grid(True)
plt.show()

```



```
In [44]: plt.figure(figsize=(12, 6))
plt.plot(y_test.values[:50], label='Actual Values', alpha=0.7)
plt.plot(y_pred_lasso[:50], label='Predicted Values (Lasso)', alpha=0.7)
plt.title('Actual vs. Predicted Values (Lasso) - First 50 Samples')
plt.xlabel('Sample Index')
plt.ylabel('Price')
plt.legend()
plt.grid(True)
plt.show()
```




```
In [45]: plt.figure(figsize=(12, 6))
```



```

plt.plot(y_test.values[:50], label='Actual Values', alpha=0.7)
plt.plot(y_pred_ridge[:50], label='Predicted Values (Ridge)', alpha=0.7)
plt.title('Actual vs. Predicted Values (Ridge) - First 50 Samples')
plt.xlabel('Sample Index')
plt.ylabel('Price')
plt.legend()
plt.grid(True)
plt.show()

```



```

In [46]: print("\n--- Model Performance Summary ---")
print("\nLinear Regression:")
print(f"  R2 Score: {r2_linear:.4f}")
print(f"  MSE: {mse_linear:.2f}")
print(f"  RMSE: {rmse_linear:.2f}")

print("\nRidge Regression:")
print(f"  R2 Score: {r2_ridge:.4f}")
print(f"  MSE: {mse_ridge:.2f}")
print(f"  RMSE: {rmse_ridge:.2f}")

print("\nLasso Regression:")
print(f"  R2 Score: {r2_lasso:.4f}")
print(f"  MSE: {mse_lasso:.2f}")
print(f"  RMSE: {rmse_lasso:.2f}")

```

--- Model Performance Summary ---

Linear Regression:

R2 Score: 0.6426

MSE: 152722264764.62

RMSE: 390796.96

Ridge Regression:

R2 Score: 0.6529

MSE: 148308986100.43

RMSE: 385109.06

Lasso Regression:

R2 Score: 0.6530

MSE: 148275760119.93

RMSE: 385065.92

In [47]: `import pickle`

```
lasso_path = "/content/drive/MyDrive/AI ML_lab_sem4/exp 11,12/Exp12_models/lasso.pkl"
with open(lasso_path, 'wb') as f:
    pickle.dump(lasso, f)

ridge_path = "/content/drive/MyDrive/AI ML_lab_sem4/exp 11,12/Exp12_models/ridge.pkl"
with open(ridge_path, 'wb') as f:
    pickle.dump(ridge, f)

scaler_path = "/content/drive/MyDrive/AI ML_lab_sem4/exp 11,12/Exp12_models/scaler.pkl"
with open(scaler_path, 'wb') as f:
    pickle.dump(scaler, f)

print("Models and scaler saved successfully!")
```

Models and scaler saved successfully!

Applied Machine Learning

School of Computer Science, UPES, Dehradun.



MINI PROJECT

B.TECH. - IV Semester

Jan - May. (2026)

SUBMITTED TO :-

Dr. Sahinur Rahman Laskar ,

Assistant Professor,

SoCS, UPES

SUBMITTED BY :-

Name : Pranav Joshi

Sap ID : 590012855

Batch : 19

Automobile Price Prediction

Dataset	automobiles.csv (UCI Automobile Dataset) Data link
Target Variable	Price (continuous numeric — USD)

1. Problem Definition

The goal of this project is to build a machine learning regression model that can accurately predict the price of an automobile given its technical and physical specifications. Unlike classification tasks, here the output is a continuous numeric value (price in USD), making it a pure regression problem.

Parameter	Detail
Target Variable (y)	price the listed selling price of the car in USD
Nature of Target	Continuous numeric; ranges from ~\$5,000 to ~\$45,000 in this dataset

2. Data Understanding


The dataset is the UCI Automobile Dataset loaded from automobiles.csv. Each row represents one car model. The dataset contains 26 columns in total, including categorical specs (make, body-style, drive-wheels) and numerical measurements (engine-size, horsepower, mpg). The target column price appears in the last column.

Feature Inventory (Input Features used in X):

Numerical Features (X): normalized-losses, wheel-base, length, width, height, curb-weight, engine-size, bore, stroke, compression-ratio, horsepower, peak-rpm, city-mpg, highway-mpg.

Categorical Features (X): make, fuel-type, aspiration, num-of-doors, body-style, drive-wheels, engine-location, engine-type, num-of-cylinders, fuel-system.

Target Variable (y): price



3. Data Pre-Processing

Raw data from automobiles.csv contains missing values encoded as '?' strings, mixed numeric/text columns, and categorical features that must be converted before any ML algorithm can process them. The preprocessing pipeline performs the following steps in strict order:

Step 3a: Replace '?' with NaN

The raw CSV uses '?' as a placeholder for missing values. Replacing with NaN allows pandas/sklearn to detect and handle them properly.

```
df.replace('?', np.nan)
```

Step 3b: Drop rows with missing Price

Price is the target variable a row without a known price cannot be used for supervised learning and is removed entirely.

```
df.dropna(subset=['price'])df['price'] = pd.to_numeric(df['price'])
```

Step 3c: Cast numerical columns to float

14 columns (engine-size, horsepower, bore, etc.) are stored as object dtype due to the '?' contamination. Explicit casting enables median imputation in the next step.

```
for col in real_num_cols: df[col] = pd.to_numeric(df[col])
```

Step 3d: Median Imputation for Numerical Missing Values

Remaining NaN values in numerical columns are filled using the median (not mean) because median is robust to outliers automobile specs like horsepower can have extreme values that would distort mean imputation.

```
imputer = SimpleImputer(strategy='median')df[real_num_cols] =  
imputer.fit_transform(df[real_num_cols])
```

Step 3e: One-Hot Encoding of Categorical Features

Categorical columns (make, body-style, fuel-type, drive-wheels, engine-type, etc.) are converted into binary (0/1) indicator columns using One-Hot Encoding. drop_first=True avoids the dummy-variable trap (multicollinearity).

```
df_encoded = pd.get_dummies(df, drop_first=True)
```

Step 3f: Feature Scaling StandardScaler

All numerical features are standardised to zero-mean and unit-variance. This is critical for regularised models (Ridge, Lasso) because they penalise coefficients — without scaling, features with large ranges dominate the penalty term.

```
scaler = StandardScaler()X_scaled = scaler.fit_transform(X)
```

4. Splitting

After preprocessing, the full feature matrix X_scaled and target vector y are split into training and testing subsets using sklearn's train_test_split.

Split Parameter	Value	Reasoning
test_size	0.2 (20%)	Standard hold-out ratio. Leaves enough data for reliable evaluation without shrinking the training set too much.
Training samples	~160 rows	80% of the cleaned dataset used to fit all regression models.

Testing samples	~40 rows	20% completely unseen during training used only for final evaluation.
-----------------	----------	---

```
X_train, X_test, y_train, y_test = train_test_split(X_scaled, y, test_size=0.2, random_state=42)
```

5. Model Training

Five regression algorithms are trained and compared. Each model was chosen for a specific reason based on the nature of the automobile dataset:

1. Linear Regression

Acts as a benchmark assuming linear price-feature relationships; no hyperparameters required for simple performance validation.

```
model_lr = LinearRegression()model_lr.fit(X_train, y_train)
```

2. Ridge Regression (L2 Regularisation)

Prevents overfitting in datasets with many dummy variables by using L2 penalties to shrink coefficients.

```
model_ridge = Ridge(alpha=1.0)model_ridge.fit(X_train, y_train)
```

3. Lasso Regression (L1 Regularisation)

Uses L1 penalties for automatic feature selection, zeroing out irrelevant features to create a sparser model.

```
model_lasso = Lasso(alpha=1.0, max_iter=10000)model_lasso.fit(X_train, y_train)
```

4. Decision Tree Regressor

Handles non-linear price impacts across feature regions without requiring data scaling for original splits.

```
model_dt = DecisionTreeRegressor(random_state=42)model_dt.fit(X_train, y_train)
```

5. Random Forest Regressor

Ensemble of 100 trees that reduces variance, captures complex feature interactions, and provides robustness against price outliers.

```
model_rf = RandomForestRegressor(random_state=42)model_rf.fit(X_train, y_train)
```


6. Model testing and evaluation



After training, every model predicts prices on the held-out test set and is evaluated using three complementary methods: numeric metrics, visual actual-vs-predicted scatter plots, and a final side-by-side comparison chart.

Evaluation Metrics Explained:

Metric	Explanation
R ² Score	A statistical measure that represents the proportion of the variance for a dependent variable that's explained by the independent variables in the model.
MSE (Mean Squared Error)	The average of the squared differences between the actual values and the predicted values.

Final Comparison :

Model	R2	MSE	RMSE
Linear regression	0.8750	14007858	3742
Ridge	0.9304	7794482	2791
Lasso	0.8831	13094871	3618
Decision tree	0.9217	8770303	2961
Random forest	0.9685	3533105	1879



7. Conclusion

- Data preprocessing was the most critical step , The dataset contained '?' as missing values and mixed data types across 26 columns, requiring type conversion, imputation, and label encoding before any model could be trained.
- Random Forest significantly outperformed all other models, achieving an R^2 of 0.9685 and an RMSE of ~1880 — far better than plain Linear Regression ($R^2 = 0.875$, RMSE ~3743), confirming that ensemble methods capture the complex, non-linear interactions in car pricing data.
- Regularization improved over basic linear regression but has limits. Ridge ($R^2 = 0.930$) outperformed both Linear and Lasso, showing that L2 regularization helps control overfitting on this high-dimensional data; however, neither could rival tree-based models.
- Feature scaling is essential for regularized models. Ridge and Lasso depend on uniform penalization across features, making StandardScaler a prerequisite; tree-based models are inherently scale-invariant yet still benefited from a consistent pipeline.



```
In [ ]: import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler, LabelEncoder
from sklearn.impute import SimpleImputer
from sklearn.linear_model import LinearRegression, Ridge, Lasso, ElasticNet, SVR
from sklearn.svm import SVC
from sklearn.neural_network import MLPClassifier
from sklearn.metrics import mean_squared_error, r2_score, classification_report
from sklearn.datasets import load_wine
from sklearn.tree import DecisionTreeRegressor
from sklearn.ensemble import RandomForestRegressor
```

```
In [ ]: df = pd.read_csv('/content/drive/MyDrive/AIML_lab_sem4/Project/automobiles.csv')
```

```
In [ ]: df.head()
```

```
Out[ ]:
```

	symboling	normalized- losses	make	fuel- type	aspiration	num- of- doors	body- style	drive- wheels
0	3	?	alfa-romero	gas	std	two	convertible	rwd
1	3	?	alfa-romero	gas	std	two	convertible	rwd
2	1	?	alfa-romero	gas	std	two	hatchback	rwd
3	2	164	audi	gas	std	four	sedan	fwd
4	2	164	audi	gas	std	four	sedan	4wd

5 rows × 26 columns

```
In [ ]: df = df.replace('?', np.nan)
```

```
In [ ]: df.info()
```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 205 entries, 0 to 204
Data columns (total 26 columns):
#   Column                Non-Null Count  Dtype
---  ---
0   symboling              205 non-null    int64
1   normalized-losses      164 non-null    object
2   make                   205 non-null    object
3   fuel-type              205 non-null    object
4   aspiration              205 non-null    object
5   num-of-doors           203 non-null    object
6   body-style             205 non-null    object
7   drive-wheels           205 non-null    object
8   engine-location        205 non-null    object
9   wheel-base             205 non-null    float64
10  length                 205 non-null    float64
11  width                  205 non-null    float64
12  height                 205 non-null    float64
13  curb-weight            205 non-null    int64
14  engine-type            205 non-null    object
15  num-of-cylinders       205 non-null    object
16  engine-size            205 non-null    int64
17  fuel-system            205 non-null    object
18  bore                   201 non-null    object
19  stroke                 201 non-null    object
20  compression-ratio      205 non-null    float64
21  horsepower             203 non-null    object
22  peak-rpm               203 non-null    object
23  city-mpg               205 non-null    int64
24  highway-mpg            205 non-null    int64
25  price                  201 non-null    object
dtypes: float64(5), int64(5), object(16)
memory usage: 41.8+ KB

```

```
In [ ]: df.shape
```

```
Out[ ]: (205, 26)
```

```
In [ ]: df = df.dropna(subset=['price'])
df['price'] = pd.to_numeric(df['price'])
```

```
In [ ]: test_case_row = df.iloc[[0]].copy()
df = df.iloc[1:].copy()
```

```
In [ ]: real_num_cols = ['normalized-losses', 'wheel-base', 'length', 'width', 'height',
                        'engine-size', 'bore', 'stroke', 'compression-ratio', 'horsepower',
                        'peak-rpm', 'city-mpg', 'highway-mpg']
```

```
In [ ]: for col in real_num_cols:
        df[col] = pd.to_numeric(df[col])
```

```
In [ ]: df.info()
```

```

<class 'pandas.core.frame.DataFrame'>
Index: 200 entries, 1 to 204
Data columns (total 26 columns):
#   Column                Non-Null Count  Dtype
---  ---
0   symboling              200 non-null    int64
1   normalized-losses      164 non-null    float64
2   make                   200 non-null    object
3   fuel-type              200 non-null    object
4   aspiration              200 non-null    object
5   num-of-doors           198 non-null    object
6   body-style             200 non-null    object
7   drive-wheels           200 non-null    object
8   engine-location        200 non-null    object
9   wheel-base            200 non-null    float64
10  length                 200 non-null    float64
11  width                  200 non-null    float64
12  height                 200 non-null    float64
13  curb-weight            200 non-null    int64
14  engine-type            200 non-null    object
15  num-of-cylinders       200 non-null    object
16  engine-size            200 non-null    int64
17  fuel-system            200 non-null    object
18  bore                   196 non-null    float64
19  stroke                 196 non-null    float64
20  compression-ratio      200 non-null    float64
21  horsepower             198 non-null    float64
22  peak-rpm               198 non-null    float64
23  city-mpg               200 non-null    int64
24  highway-mpg            200 non-null    int64
25  price                  200 non-null    int64
dtypes: float64(10), int64(6), object(10)
memory usage: 42.2+ KB

```

```

In [ ]: imputer = SimpleImputer(strategy='median')
df[real_num_cols] = imputer.fit_transform(df[real_num_cols])

```


Exploratory Data Analysis (EDA)

Let's start by visualizing the distribution of the target variable, `price`.

```

In [ ]: plt.figure(figsize=(10, 6))
sns.histplot(df['price'], bins=30, kde=True)
plt.title('Distribution of Car Price')
plt.xlabel('Price')
plt.ylabel('Frequency')
plt.show()

```



Now, let's look at the distributions of some other key numerical features like 'horsepower', 'engine-size', and 'city-mpg'.


```
In [ ]: plt.figure(figsize=(15, 5))

plt.subplot(1, 3, 1)
sns.histplot(df['horsepower'], bins=20, kde=True)
plt.title('Distribution of Horsepower')


plt.subplot(1, 3, 2)
sns.histplot(df['engine-size'], bins=20, kde=True)
plt.title('Distribution of Engine Size')

plt.subplot(1, 3, 3)
sns.histplot(df['city-mpg'], bins=20, kde=True)
plt.title('Distribution of City MPG')

plt.tight_layout()
plt.show()
```



```
In [ ]: plt.figure(figsize=(12, 5))
sns.boxplot(x='make', y='price', data=df)
plt.xticks(rotation=90)
plt.title('Price by Car Make')
plt.xlabel('Make')
plt.ylabel('Price')
plt.tight_layout()
plt.show()
```



```
In [ ]: df_encoded = pd.get_dummies(df, drop_first=True)
```

```
In [ ]: X = df_encoded.drop('price', axis=1)
y = df_encoded['price']
```

```
In [ ]: scaler = StandardScaler()
X_scaled = scaler.fit_transform(X)
```

```
In [ ]: X_train, X_test, y_train, y_test = train_test_split(X_scaled, y, test_size=0.2)
```

```
In [ ]: results = []
```



```
print("Preprocessing complete. Training set size:", X_train.shape[0])
```

Preprocessing complete. Training set size: 160

Linear regression

```
In [ ]: model_lr = LinearRegression()
```

```
In [ ]: model_lr.fit(X_train, y_train)
```

```
Out[ ]: > LinearRegression
```

```
LinearRegression()
```


```
In [ ]: preds_lr = model_lr.predict(X_test)
```

```
In [ ]: mse_lr = mean_squared_error(y_test, preds_lr)
r2_lr = r2_score(y_test, preds_lr)
rmse_lr = np.sqrt(mse_lr)
```

```
In [ ]: results.append(("Model": "Linear", "MSE": mse_lr, "R2": r2_lr, "RMSE": rmse_lr))
print(f"Linear Regression ->\n R2: {r2_lr:.4f},\n MSE: {mse_lr:.2f},\n RMSE: {
```

Linear Regression ->
R2: 0.8750,
MSE: 14007858.67,
RMSE: 3742.71

```
In [ ]: plt.scatter(y_test, preds_lr, alpha=0.7, color='blue')
plt.plot([y_test.min(), y_test.max()], [y_test.min(), y_test.max()], 'r--', lw=2)
plt.title("Linear Regression: Predicted vs Actual")
plt.xlabel("Actual Price")
plt.ylabel("Predicted Price")
plt.show()
```




```
In [ ]: model_ridge = Ridge(alpha=1.0)
model_ridge.fit(X_train, y_train)
preds_ridge = model_ridge.predict(X_test)

In [ ]: mse_ridge = mean_squared_error(y_test, preds_ridge)
r2_ridge = r2_score(y_test, preds_ridge)
rmse_ridge = np.sqrt(mse_ridge)

In [ ]: results.append({"Model": "Ridge", "MSE": mse_ridge, "R2": r2_ridge, "RMSE": rmse_ridge})

In [ ]: print(f"Ridge Regression -> R2: {r2_ridge:.4f}, MSE: {mse_ridge:.2f}, RMSE: {rmse_ridge:.2f}")
Ridge Regression -> R2: 0.9384, MSE: 7794482.02, RMSE: 2791.86

In [ ]: plt.scatter(y_test, preds_ridge, alpha=0.7, color='green')
plt.plot([y_test.min(), y_test.max()], [y_test.min(), y_test.max()], 'r--', lw=2)
plt.title("Ridge Regression: Predicted vs Actual")
plt.xlabel("Actual Price")
plt.ylabel("Predicted Price")
plt.show()
```




```
In [ ]: model_lasso = Lasso(alpha=1.0, max_iter=10000)
model_lasso.fit(X_train, y_train)
preds_lasso = model_lasso.predict(X_test)

In [ ]: mse_lasso = mean_squared_error(y_test, preds_lasso)
r2_lasso = r2_score(y_test, preds_lasso)
rmse_lasso = np.sqrt(mse_lasso)

In [ ]: results.append({"Model": "Lasso", "MSE": mse_lasso, "R2": r2_lasso, "RMSE": rmse_lasso})

In [ ]: print(f"Lasso Regression -> R2: {r2_lasso:.4f}, MSE: {mse_lasso:.2f}, RMSE: {rmse_lasso:.2f}")
Lasso Regression -> R2: 0.8831, MSE: 13894871.84, RMSE: 3618.68

In [ ]: plt.scatter(y_test, preds_lasso, alpha=0.7, color='purple')
plt.plot([y_test.min(), y_test.max()], [y_test.min(), y_test.max()], 'r--', lw=2)
plt.title("Lasso Regression: Predicted vs Actual")
plt.xlabel("Actual Price")
plt.ylabel("Predicted Price")
plt.show()
```




Decision Tree Regressor

```
In [ ]: model_dt = DecisionTreeRegressor(random_state=42)
model_dt.fit(X_train, y_train)
preds_dt = model_dt.predict(X_test)
```

```
In [ ]: mse_dt = mean_squared_error(y_test, preds_dt)
r2_dt = r2_score(y_test, preds_dt)
rmse_dt = np.sqrt(mse_dt)
results.append({"Model": "Decision Tree", "MSE": mse_dt, "R2": r2_dt, "RMSE": rmse_dt})
print(f"Decision Tree Regressor -> R2: {r2_dt:.4f}, MSE: {mse_dt:.2f}, RMSE: {rmse_dt:.2f}")
```

Decision Tree Regressor -> R2: 0.9217, MSE: 8770303.55, RMSE: 2961.47

```
In [ ]: plt.scatter(y_test, preds_dt, alpha=0.7, color='orange')
plt.plot([y_test.min(), y_test.max()], [y_test.min(), y_test.max()], 'r--', lw=2)
plt.title("Decision Tree Regression: Predicted vs Actual")
plt.xlabel("Actual Price")
plt.ylabel("Predicted Price")
plt.show()
```




Random Forest Regressor

```
In [ ]: model_rf = RandomForestRegressor(random_state=42)
model_rf.fit(X_train, y_train)
preds_rf = model_rf.predict(X_test)
```

```
In [ ]: mse_rf = mean_squared_error(y_test, preds_rf)
r2_rf = r2_score(y_test, preds_rf)
rmse_rf = np.sqrt(mse_rf)
results.append({"Model": "Random Forest", "MSE": mse_rf, "R2": r2_rf, "RMSE":
print(f"Random Forest Regressor -> R2: {r2_rf:.4f}, MSE: {mse_rf:.2f}, RMSE: {
```

Random Forest Regressor -> R2: 0.9685, MSE: 3533105.77, RMSE: 1879.66

```
In [ ]: plt.scatter(y_test, preds_rf, alpha=0.7, color='brown')
plt.plot([y_test.min(), y_test.max()], [y_test.min(), y_test.max()], 'r--', lw
plt.title("Random Forest Regression: Predicted vs Actual")
plt.xlabel("Actual Price")
plt.ylabel("Predicted Price")
plt.show()
```



Final Model Comparison

```
In [ ]: print("\nFinal Model Comparison Summary:")
display(comparison_df)
```

Final Model Comparison Summary:

	Model	MSE	R2	RMSE
0	Linear	1.400786e+07	0.874984	3742.707399
1	Ridge	7.794482e+06	0.930436	2791.859958
2	Lasso	1.309487e+07	0.883132	3618.683605
3	Decision Tree	8.770304e+06	0.921727	2961.469829
4	Random Forest	3.533106e+06	0.968468	1879.655759

```
In [ ]: comparison_df = pd.DataFrame(results)

plt.figure(figsize=(10, 7))
sns.barplot(x='Model', y='R2', data=comparison_df, palette='magma')
plt.title("Comparison of R2 Scores (Higher is Better)")
plt.ylim(0, 1)
plt.tight_layout()
```




```
plt.show()
```

/tmp/ipykernel_6578/3386395981.py:4: FutureWarning:

Passing 'palette' without assigning 'hue' is deprecated and will be removed in v0.14.0. Assign the 'x' variable to 'hue' and set 'legend=False' for the same effect.

```
sns.barplot(x='Model', y='R2', data=comparison_df, palette='magma')
```




```
In [ ]: plt.figure(figsize=(10, 7))
sns.barplot(x='Model', y='MSE', data=comparison_df, palette='viridis')
plt.title("Comparison of MSE (Lower is Better)")
plt.tight_layout()
plt.show()
```

/tmp/ipykernel_6578/3137539581.py:2: FutureWarning:

Passing 'palette' without assigning 'hue' is deprecated and will be removed in v0.14.0. Assign the 'x' variable to 'hue' and set 'legend=False' for the same effect.

```
sns.barplot(x='Model', y='MSE', data=comparison_df, palette='viridis')
```



Health Risk Level Prediction

Dataset	Health_Risk_Dataset.csv
Target Variable	Risk_Level (4 classes: Low, Medium, High, Critical)

1. Problem Definition

The goal of this project is to build a machine learning model that can classify a patient's health risk level based on their vital signs and clinical indicators. The output is one of four discrete categories Low, Medium, High, or Critical making this a multi-class classification problem, not a regression problem.

Parameter	Detail
Target Variable (y)	Risk_Level — 4 classes: Low / Medium / High / Critical
Nature of Target	Categorical, ordinal; encoded as integers via LabelEncoder (0, 1, 2, 3)

2. Data Understanding

The dataset is loaded from Health_Risk_Dataset.csv. Each row represents one patient record. The dataset includes a non-informative ID column, one mixed-type categorical column (Consciousness), several numerical vital sign measurements, and the target column Risk_Level. A correlation matrix is computed to understand feature relationships before modelling begins.

Feature Inventory (Input Features used in X):

Numerical Features (X):

- Respiratory_Rate
- Oxygen_Saturation
- O2_Scale
- Systolic_BP
- Heart_Rate
- Temperature
- On_Oxygen



Categorical Features (X):

- Consciousness — Alert (A) / Responds to Pain (P) / Unresponsive (U)

Target Variable (y):

- Risk_Level — Normal/Low / Medium / High

Class Distribution :



3. Data Pre-Processing

The raw dataset requires three targeted preprocessing steps before any model can be trained.

Step 3a: Drop Patient_ID Column

Patient_ID is a unique row identifier with no predictive meaning. Including it would cause the model to memorise IDs instead of learning clinical patterns a classic form of data leakage.

```
df.drop(columns=['Patient_ID'])
```

Step 3b: Label Encode 'Consciousness'

Consciousness has 3 ordinal states (Alert < Confused < Unresponsive). LabelEncoder maps these to 0, 1, 2, preserving the inherent clinical order. One-Hot Encoding is avoided here because the ordinal relationship is medically meaningful.

```
le_cons = LabelEncoder()df['Consciousness'] = le_cons.fit_transform(df['Consciousness'])
```

Step 3c: Label Encode Target

The 4-class target (Low, Medium, High, Critical) must be integer-encoded for sklearn classifiers. le_target.classes_ is saved so original class names can be restored in classification reports and confusion matrix labels.

```
le_target = LabelEncoder()(df['Risk_Level']) = le_target.fit_transform(df['Risk_Level'])
```

Step 3d: Feature Scaling

Vital signs vary in magnitude; without standardization, larger values dominate Logistic Regression. Fit your scaler only on X_train to prevent data leakage and ensure model integrity.

```
scaler = StandardScaler()X_train_scaled = scaler.fit_transform(X_train)X_test_scaled = scaler.transform(X_test)
```

4. Splitting Point

The preprocessed feature matrix X and target vector y are split before scaling is applied, ensuring the scaler learns only from training statistics.

Split Parameter	Value	Reasoning
test_size	0.2 (20%)	Standard 80/20 split; enough test samples for reliable metric estimation across all 4 risk classes.
Train size	800 rows	Used to fit all 5 models. Scaler is also fitted only on this subset.
Test size	200 rows	Completely held-out. Used only for predict() and metric calculation.

5. Model Training

Five classification models are trained, spanning from simple linear classifiers to ensemble methods. Each was chosen for a specific reason tied to the multi-class nature of health risk prediction:

1. Logistic Regression

This unregularized multinomial baseline uses Softmax to model four classes simultaneously. It serves as a benchmark; if regularization provides no improvement, the features are already well-separated.

```
m1 = LogisticRegression( penalty=None, multi_class='multinomial', max_iter=10000)m1.fit(X_train_scaled, y_train)
```

2. Logistic Regression(L1)

Using a SAGA solver, this L1-regularized model (C=0.1) performs feature selection by zeroing out redundant coefficients. This is ideal for health data where overlapping vital signs often provide duplicate information. .

```
m2 = LogisticRegression( penalty='l1', solver='saga', multi_class='multinomial', max_iter=5000, C=0.1)
```

3. Logistic Regression (L2)

This **L2-regularized** model (C=0.01) shrinks coefficients toward zero without eliminating them. It is used to determine if keeping all features at small weights outperforms the selection-heavy L1 approach.

```
m3 = LogisticRegression( penalty='l2', solver='saga', multi_class='multinomial', max_iter=1000, C=0.01)
```

4. Decision Tree Classifier

This rule-based classifier uses a depth of five to ensure medical interpretability and prevent overfitting, mirroring clinical diagnostic logic.

```
m4 = DecisionTreeClassifier( max_depth=5, random_state=42)
```

5. Bagging Classifier

By averaging 50 diverse trees, this ensemble reduces variance and stabilizes predictions, balancing computational speed with robust, reliable accuracy.

```
m5 = BaggingClassifier( estimator=DecisionTreeClassifier(max_depth=5), n_estimators=50, random_state=42)
```

6. XG Boost

This gradient boosting model builds trees sequentially, with each new tree correcting errors from previous ones. It offers high performance and efficiency through advanced regularization and optimized computation.

```
XGBClassifier(random_state=RANDOM_STATE, use_label_encoder=False')
```


6. Test Model


Each model is evaluated on the held-out test set using four complementary methods: a full classification report, per-class confusion matrices, a comparative bar chart, and macro-averaged ROC curves with AUC scores.

Evaluation Metrics Explained:

Metric	Explanation
Precision (macro avg)	Macro average precision measures correct class predictions, emphasizing low false positives.
Recall (macro avg)	Recall measures the fraction of true cases correctly identified, which is vital for catching all critical patients.
F1-Score (macro avg)	The harmonic mean of precision and recall, balancing both metrics. It is used here to ensure the model doesn't ignore minority risk classes.
ROC + AUC	evaluates each class binary, averaging the results. Using probability scores, it compares all models to determine discrimination excellence.

	Model	Precision	Recall	F1-Score	AUC	Accuracy
1	XGBoost	0.961230	0.956258	0.958598	0.998877	0.955
2	Logistic Regression	0.961400	0.950071	0.954576	0.996385	0.955
3	Logit L1	0.945012	0.937838	0.940833	0.995876	0.940
4	Logit L2	0.925307	0.897222	0.907773	0.980754	0.910
5	Bagging DT	0.914488	0.899364	0.905934	0.987020	0.900
6	Decision Tree	0.906018	0.891605	0.897761	0.961525	0.890





7. Conclusion

1. Stratified splitting ensured fair class representation. With 4 imbalanced classes , Medium (306), High (279), Low (255), and Normal (160) , using stratify=y preserved each class's proportion in the test set, making macro F1 and AUC scores reliable across all risk groups, especially the minority Normal class.
2. XGBoost emerged as the best overall model, achieving the highest macro F1 (0.959) and AUC (0.999), along with 95.5% accuracy , demonstrating that gradient boosting captures complex, non-linear interactions between vital signs far better than linear or shallow tree-based models.
3. Logistic Regression variants revealed clear regularisation tradeoffs. Vanilla LR (F1 = 0.955) outperformed L1 (0.941) and L2 (0.908), showing that aggressive regularisation hurt performance here ,L2 at C=0.01 notably degraded recall for Low and Normal classes by over-shrinking informative feature weights.
4. Bagging improved on a single Decision Tree but couldn't close the gap with linear models. The single Decision Tree (F1 = 0.898) was the weakest model; 50 bootstrapped trees in the BaggingClassifier raised the macro F1 to 0.906, confirming variance reduction , but both lagged behind even L2 Logistic Regression.
5. High AUC across all models confirmed strong probabilistic separation of risk levels. Every model scored above 0.96 AUC, meaning all models rank patients reliably by risk , making threshold tuning at deployment a viable strategy to further optimise recall for the Critical and High risk groups in a hospital triage context.



```
In [257]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
import warnings
warnings.filterwarnings('ignore')

from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler, LabelEncoder, Label_Binarizer
from sklearn.linear_model import LogisticRegression
from sklearn.tree import DecisionTreeClassifier
from sklearn.ensemble import BaggingClassifier
from sklearn.metrics import (confusion_matrix, ConfusionMatrixDisplay, classification_score, precision_score, recall_score,
                             roc_curve, auc)

from xgboost import XGBClassifier
```

```
In [258]: df = pd.read_csv('/content/drive/MyDrive/AIML_lab_sem4/Project/Health_Risk_Data.csv')
```

```
In [259]: df.head()
```

```
Out[259]:
```

	Patient_ID	Respiratory_Rate	Oxygen_Saturation	O2_Scale	Systolic_BP	Heart_Rate
0	P0522	25	96	1	97	78
1	P0738	28	92	2	116	85
2	P0741	29	91	1	79	82
3	P0661	24	96	1	95	75
4	P0412	20	96	1	97	72

```
In [260]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1000 entries, 0 to 999
Data columns (total 10 columns):
#   Column                Non-Null Count  Dtype  
---  -
0   Patient_ID            1000 non-null  object 
1   Respiratory_Rate      1000 non-null  int64  
2   Oxygen_Saturation     1000 non-null  int64  
3   O2_Scale              1000 non-null  int64  
4   Systolic_BP           1000 non-null  int64  
5   Heart_Rate            1000 non-null  int64  
6   Temperature           1000 non-null  float64 
7   Consciousness         1000 non-null  object 
8   On_Oxygen             1000 non-null  int64  
9   Risk_Level            1000 non-null  object 
dtypes: float64(1), int64(6), object(3)
memory usage: 78.3+ KB
```

```
In [261]: df.isnull().sum()
```



```
Out[261]:
```

Patient_ID	0
Respiratory_Rate	0
Oxygen_Saturation	0
O2_Scale	0
Systolic_BP	0
Heart_Rate	0
Temperature	0
Consciousness	0
On_Oxygen	0
Risk_Level	0

dtype: int64

```
In [262]: df.describe()
```

```
Out[262]:
```

	Respiratory_Rate	Oxygen_Saturation	O2_Scale	Systolic_BP	Heart_R
count	1000.000000	1000.000000	1000.000000	1000.000000	1000.000
mean	21.511000	92.590000	1.124000	106.160000	98.460
std	5.287517	4.473020	0.329746	17.897562	19.694
min	12.000000	74.000000	1.000000	50.000000	60.000
25%	17.000000	90.000000	1.000000	94.000000	84.000
50%	20.000000	94.000000	1.000000	109.000000	95.500
75%	25.000000	96.000000	1.000000	119.000000	109.000
max	40.000000	100.000000	2.000000	146.000000	163.000

```
In [263]: TARGET = 'Risk_Level'
```

```
print(f"Target column : '{TARGET}'")
print('\nClass Distribution (before encoding):')
print(df[TARGET].value_counts())
print(f'\nTotal classes : {df[TARGET].nunique()}')
```

Target column : 'Risk_Level'

Class Distribution (before encoding):

```
Risk_Level
Medium    306
High      279
Low       255
Normal    160
Name: count, dtype: int64
```


Total classes : 4

```
In [264]: print('Generating EDA Plots...')

counts = df[TARGET].value_counts().sort_index()
colors_cls = ['#4C72B0', '#DD8452', '#55A868', '#C44E52', '#B17283']

plt.figure(figsize=(7, 4))
bars = plt.bar(counts.index.astype(str), counts.values,
               color=colors_cls[:len(counts)], edgecolor='black')
for bar in bars:
    plt.text(bar.get_x() + bar.get_width() / 2,
             bar.get_height() + max(counts.values) * 0.01,
             str(bar.get_height()), ha='center', fontweight='bold', fontsize=10)
plt.title(f'Class Distribution - {TARGET}', fontsize=13, fontweight='bold')
plt.xlabel('Risk Level')
plt.ylabel('Count')
plt.tight_layout()
plt.savefig('plot1_class_distribution.png', dpi=150)
plt.show()
print('Saved: plot1_class_distribution.png')
```


Generating EDA Plots...



Saved: plot1_class_distribution.png

```
In [265]: numeric_df = df.select_dtypes(include=[np.number])

plt.figure(figsize=(12, 8))
sns.heatmap(numeric_df.corr(), annot=True, fmt='.2f',
            cmap='coolwarm', linewidths=0.5, square=True)
plt.title('Correlation Heatmap (Numeric Features)', fontsize=13, fontweight='b')
plt.tight_layout()
plt.savefig('plot3_correlation_heatmap.png', dpi=150)
plt.show()
print('Saved: plot3_correlation_heatmap.png')
```



Saved: plot3_correlation_heatmap.png


```
In [266]: KEY_FEATURE = numeric_df.columns[0] # change to preferred column if desired

plt.figure(figsize=(8, 4))
for cls, clr in zip(df[TARGET].unique(), colors_cls):
    subset = df[df[TARGET] == cls][KEY_FEATURE].dropna()
    plt.hist(subset, bins=35, alpha=0.55, color=clr,
```

```

        label=str(cls), edgecolor='black')
plt.title(f'{KEY_FEATURE} Distribution by Risk Level', fontsize=12, fontweight
plt.xlabel(KEY_FEATURE)
plt.ylabel('Count')
plt.legend(title='Risk Level')
plt.tight_layout()
plt.savefig('plot4_feature_by_class.png', dpi=150)
plt.show()
print('Saved: plot4_feature_by_class.png')

```



Saved: plot4_feature_by_class.png

Preprocessing

```
In [267]: df = df.drop(columns=['Patient_ID'], errors='ignore')
```

```
In [268]: le_cons = LabelEncoder()
df['Consciousness'] = le_cons.fit_transform(df['Consciousness'])
print('Encoded: Consciousness')
```

Encoded: Consciousness

```
In [269]: le_target = LabelEncoder()
df[TARGET] = le_target.fit_transform(df[TARGET])
```

```
In [270]: X = df.drop(columns=[TARGET])
y = df[TARGET]
```

```
In [271]: RANDOM_STATE = 42

X_train, X_test, y_train, y_test = train_test_split(
    X, y, test_size=0.2, random_state=RANDOM_STATE, stratify=y
)
```

```

In [272]: scaler = StandardScaler()
X_train_scaled = scaler.fit_transform(X_train)
X_test_scaled = scaler.transform(X_test)

In [273]: m1 = LogisticRegression(penalty=None, multi_class='multinomial',
                                solver='lbfgs', max_iter=10000, random_state=RANDOM_STATE)
m1.fit(X_train_scaled, y_train)
p1 = m1.predict(X_test_scaled)

print('--- Model 1: Logistic Regression ---')
print(classification_report(y_test, p1, target_names=le_target.classes_))

--- Model 1: Logistic Regression ---
              precision    recall  f1-score   support

      High         1.00        0.95        0.97         56
       Low         0.91        0.98        0.94         51
     Medium         0.94        0.97        0.95         61
      Normal         1.00        0.91        0.95         32

 accuracy          0.96
 macro avg          0.96
 weighted avg       0.96

In [274]: print('Manually tuning L1 Logistic Regression...')
m2 = LogisticRegression(penalty='l1', solver='saga', multi_class='multinomial',
                        max_iter=5000, C=0.1, random_state=RANDOM_STATE)
m2.fit(X_train_scaled, y_train)
p2 = m2.predict(X_test_scaled)
print('L1 Logistic Regression (C=0.1) fitted.')

print('\n--- Model 2: Logistic Regression L1 ---')
print(classification_report(y_test, p2, target_names=le_target.classes_))

Manually tuning L1 Logistic Regression...
L1 Logistic Regression (C=0.1) fitted.

--- Model 2: Logistic Regression L1 ---
              precision    recall  f1-score   support

      High         1.00        0.96        0.98         56
       Low         0.92        0.88        0.90         51
     Medium         0.89        0.97        0.93         61
      Normal         0.97        0.94        0.95         32

 accuracy          0.94
 macro avg          0.95
 weighted avg       0.94

In [275]: print('Manually tuning L2 Logistic Regression...')
m3 = LogisticRegression(penalty='l2', solver='lbfgs', multi_class='multinomial',
                        max_iter=1000, C=0.01, random_state=RANDOM_STATE)

```



```

m3.fit(X_train_scaled, y_train)
p3 = m3.predict(X_test_scaled)
print('L2 Logistic Regression (C=0.01) fitted.')

print('\n--- Model 3: Logistic Regression L2 ---')
print(classification_report(y_test, p3, target_names=le_target.classes_))

```

Manually tuning L2 Logistic Regression...

L2 Logistic Regression (C=0.01) fitted.

```

--- Model 3: Logistic Regression L2 ---

```

	precision	recall	f1-score	support
High	1.00	0.95	0.97	56
Low	0.85	0.86	0.85	51
Medium	0.86	0.97	0.91	61
Normal	1.00	0.81	0.90	32
accuracy			0.91	200
macro avg	0.93	0.90	0.91	200
weighted avg	0.92	0.91	0.91	200

```

In [276]: m4 = DecisionTreeClassifier(max_depth=5, random_state=RANDOM_STATE)
m4.fit(X_train_scaled, y_train)
p4 = m4.predict(X_test_scaled)

print('--- Model 4: Decision Tree ---')
print(classification_report(y_test, p4, target_names=le_target.classes_))

```

```

--- Model 4: Decision Tree ---

```

	precision	recall	f1-score	support
High	0.98	0.93	0.95	56
Low	0.81	0.86	0.84	51
Medium	0.83	0.87	0.85	61
Normal	1.00	0.91	0.95	32
accuracy			0.89	200
macro avg	0.91	0.89	0.90	200
weighted avg	0.90	0.89	0.89	200

```

In [277]: m5 = BaggingClassifier(estimator=DecisionTreeClassifier(max_depth=5),
                                n_estimators=50, random_state=RANDOM_STATE)
m5.fit(X_train_scaled, y_train)
p5 = m5.predict(X_test_scaled)

print('--- Model 5: Bagging Classifier ---')
print(classification_report(y_test, p5, target_names=le_target.classes_))

```

```

--- Model 5: Bagging Classifier ---
precision    recall  f1-score   support

   High      0.98      0.95      0.96        56
   Low       0.84      0.84      0.84        51
  Medium     0.83      0.90      0.87        61
   Normal    1.00      0.91      0.95        32

 accuracy    0.90        200
 macro avg   0.91      0.90      0.91        200
weighted avg   0.90      0.90      0.90        200

```

```

In [278]: m_xgb = XGBClassifier(random_state=RANDOM_STATE, use_label_encoder=False, eval
m_xgb.fit(X_train_scaled, y_train)
p_xgb = m_xgb.predict(X_test_scaled)

print('\n--- Model 6: XGBoost Classifier ---')
print(classification_report(y_test, p_xgb, target_names=le_target.classes_))

```

```

--- Model 6: XGBoost Classifier ---
precision    recall  f1-score   support

   High      0.96      0.96      0.96        56
   Low       0.96      0.94      0.95        51
  Medium     0.92      0.95      0.94        61
   Normal    1.00      0.97      0.98        32

 accuracy    0.95        200
 macro avg   0.96      0.96      0.96        200
weighted avg   0.96      0.95      0.96        200

```


```

In [280]: bar_colors_f1 = [
'green'      if v == max(f1_macro_scores) else
'red'        if v == min(f1_macro_scores) else
'steelblue'
for v in f1_macro_scores
]

plt.figure(figsize=(12, 5))
bars = plt.bar(model_names, f1_macro_scores,
               color=bar_colors_f1, edgecolor='black', width=0.5)
for bar in bars:
    plt.text(bar.get_x() + bar.get_width() / 2,
             bar.get_height() + 0.005,
             f'{bar.get_height():.4f}',
             ha='center', fontsize=9, fontweight='bold')
plt.title('F1-Score (Macro) - All Models\n(Green = Best | Red = Worst)',
          fontsize=13, fontweight='bold')
plt.ylabel('Macro F1-Score')
plt.ylim(0, 1.15)
plt.xticks(rotation=45, ha='right')
plt.grid(axis='y', linestyle='--', alpha=0.6)

```

```
plt.tight_layout()
plt.savefig('plot5_f1_bar.png', dpi=150)
plt.show()
print('Saved: plot5_f1_bar.png')
```




Saved: plot5_f1_bar.png

```
In [281]: x = np.arange(len(model_names))
w = 0.2


plt.figure(figsize=(15, 5))
plt.bar(x - w, precision_macro_scores, w,
        label='Precision', color='steelblue', edgecolor='black')
plt.bar(x, recall_macro_scores, w,
        label='Recall', color='coral', edgecolor='black')
plt.bar(x + w, f1_macro_scores, w,
        label='F1-Score', color='green', edgecolor='black')

plt.xticks(x, model_names, rotation=45, ha='right', fontsize=9)
plt.ylabel('Macro Average Score')
plt.ylim(0, 1.2)
plt.title('Precision | Recall | F1-Score — All Models (Grouped)',
          fontsize=13, fontweight='bold')
plt.legend()
plt.grid(axis='y', linestyle='--', alpha=0.6)
plt.tight_layout()
plt.savefig('plot6_grouped_bar.png', dpi=150)
plt.show()
print('Saved: plot6_grouped_bar.png')
```


Saved: plot6_grouped_bar.png

```
In [282]: # Model 1: Baseline
plt.figure(figsize=(6, 5))
sns.heatmap(confusion_matrix(y_test, p1), annot=True, fmt='d', cmap='Blues',
            xticklabels=le_target.classes_, yticklabels=le_target.classes_)
plt.title(f'Confusion Matrix: {model_names[0]}\nF1 = {f1_macro_scores[0]:.3f}')
plt.xlabel('Predicted')
plt.ylabel('True')
plt.show()
```




```
In [283]: # Model 2: Lasso (L1)
```


```
plt.figure(figsize=(6, 5))
sns.heatmap(confusion_matrix(y_test, p2), annot=True, fmt='d', cmap='Greens',
            xticklabels=le_target.classes_, yticklabels=le_target.classes_)
plt.title(f'Confusion Matrix: {model_names[1]}\nF1 = {f1_macro_scores[1]:.3f}')
plt.xlabel('Predicted')
plt.ylabel('True')
plt.show()
```




```
In [284]: # Model 3: Ridge (L2)
plt.figure(figsize=(6, 5))
sns.heatmap(confusion_matrix(y_test, p3), annot=True, fmt='d', cmap='Oranges',
            xticklabels=le_target.classes_, yticklabels=le_target.classes_)
plt.title(f'Confusion Matrix: {model_names[2]}\nF1 = {f1_macro_scores[2]:.3f}')
plt.xlabel('Predicted')
plt.ylabel('True')
plt.show()
```




```
In [285]: # Model 4: Decision Tree
plt.figure(figsize=(6, 5))
sns.heatmap(confusion_matrix(y_test, p4), annot=True, fmt='d', cmap='Purples',
            xticklabels=le_target.classes_, yticklabels=le_target.classes_)
plt.title(f'Confusion Matrix: {model_names[3]}\nF1 = {f1_macro_scores[3]:.3f}')
plt.xlabel('Predicted')
plt.ylabel('True')
plt.show()
```



```
In [286]: # Model 5: Random Forest
plt.figure(figsize=(6, 5))
sns.heatmap(confusion_matrix(y_test, p5), annot=True, fmt='d', cmap='Reds',
            xticklabels=le_target.classes_, yticklabels=le_target.classes_)
plt.title(f'Confusion Matrix: {model_names[4]}\nF1 = {f1_macro_scores[4]:.3f}')
plt.xlabel('Predicted')
plt.ylabel('True')
plt.show()
```



```
In [287]: # Model 6: XGBoost
plt.figure(figsize=(6, 5))
sns.heatmap(confusion_matrix(y_test, p_xgb), annot=True, fmt='d', cmap='Blues',
            xticklabels=le_target.classes_, yticklabels=le_target.classes_)
plt.title(f'Confusion Matrix: {model_names[5]}\nF1 = {f1_macro_scores[5]:.3f}')
plt.xlabel('Predicted')
plt.ylabel('True')
plt.show()
```



```
In [288]: y_test_bin = label_binarize(y_test, classes=np.unique(y_train))
n_classes = y_test_bin.shape[1]

models_roc = [
    ('Logistic Regression', m1),
    ('Logit L1',          m2),
    ('Logit L2',          m3),
    ('Decision Tree',     m4),
    ('Bagging DT',        m5),
    ('XGBoost',           m_xgb)
]

roc_colors = plt.cm.tab10(np.linspace(0, 1, len(models_roc)))
auc_scores_list = []

plt.figure(figsize=(10, 7))

y_score_m1 = m1.predict_proba(X_test_scaled)
fpr_m1, tpr_m1, roc_auc_m1 = {}, {}, {}
for i in range(n_classes):
    fpr_m1[i], tpr_m1[i], _ = roc_curve(y_test_bin[:, i], y_score_m1[:, i])
    roc_auc_m1[i] = auc(fpr_m1[i], tpr_m1[i])
all_fpr_m1 = np.unique(np.concatenate([fpr_m1[i] for i in range(n_classes)]))
mean_tpr_m1 = np.zeros_like(all_fpr_m1)
```



```

for i in range(n_classes):
    mean_tpr_m1 += np.interp(all_fpr_m1, fpr_m1[i], tpr_m1[i])
mean_tpr_m1 /= n_classes
macro_auc_m1 = auc(all_fpr_m1, mean_tpr_m1)
plt.plot(all_fpr_m1, mean_tpr_m1, color=roc_colors[0], linewidth=2,
         label=f'Logistic Regression (AUC = {macro_auc_m1:.2f})')
auc_scores_list.append({'Model': 'Logistic Regression', 'AUC': macro_auc_m1})

# --- Logit L1 ---
y_score_m2 = m2.predict_proba(X_test_scaled)
fpr_m2, tpr_m2, roc_auc_m2 = {}, {}, {}
for i in range(n_classes):
    fpr_m2[i], tpr_m2[i], _ = roc_curve(y_test_bin[:, i], y_score_m2[:, i])
    roc_auc_m2[i] = auc(fpr_m2[i], tpr_m2[i])
all_fpr_m2 = np.unique(np.concatenate([fpr_m2[i] for i in range(n_classes)]))
mean_tpr_m2 = np.zeros_like(all_fpr_m2)
for i in range(n_classes):
    mean_tpr_m2 += np.interp(all_fpr_m2, fpr_m2[i], tpr_m2[i])
mean_tpr_m2 /= n_classes
macro_auc_m2 = auc(all_fpr_m2, mean_tpr_m2)
plt.plot(all_fpr_m2, mean_tpr_m2, color=roc_colors[1], linewidth=2,
         label=f'Logit L1 (AUC = {macro_auc_m2:.2f})')
auc_scores_list.append({'Model': 'Logit L1', 'AUC': macro_auc_m2})

# --- Logit L2 ---
y_score_m3 = m3.predict_proba(X_test_scaled)
fpr_m3, tpr_m3, roc_auc_m3 = {}, {}, {}
for i in range(n_classes):
    fpr_m3[i], tpr_m3[i], _ = roc_curve(y_test_bin[:, i], y_score_m3[:, i])
    roc_auc_m3[i] = auc(fpr_m3[i], tpr_m3[i])
all_fpr_m3 = np.unique(np.concatenate([fpr_m3[i] for i in range(n_classes)]))
mean_tpr_m3 = np.zeros_like(all_fpr_m3)
for i in range(n_classes):
    mean_tpr_m3 += np.interp(all_fpr_m3, fpr_m3[i], tpr_m3[i])
mean_tpr_m3 /= n_classes
macro_auc_m3 = auc(all_fpr_m3, mean_tpr_m3)
plt.plot(all_fpr_m3, mean_tpr_m3, color=roc_colors[2], linewidth=2,
         label=f'Logit L2 (AUC = {macro_auc_m3:.2f})')
auc_scores_list.append({'Model': 'Logit L2', 'AUC': macro_auc_m3})

# --- Decision Tree ---
y_score_m4 = m4.predict_proba(X_test_scaled)
fpr_m4, tpr_m4, roc_auc_m4 = {}, {}, {}
for i in range(n_classes):
    fpr_m4[i], tpr_m4[i], _ = roc_curve(y_test_bin[:, i], y_score_m4[:, i])
    roc_auc_m4[i] = auc(fpr_m4[i], tpr_m4[i])
all_fpr_m4 = np.unique(np.concatenate([fpr_m4[i] for i in range(n_classes)]))
mean_tpr_m4 = np.zeros_like(all_fpr_m4)
for i in range(n_classes):
    mean_tpr_m4 += np.interp(all_fpr_m4, fpr_m4[i], tpr_m4[i])
mean_tpr_m4 /= n_classes
macro_auc_m4 = auc(all_fpr_m4, mean_tpr_m4)
plt.plot(all_fpr_m4, mean_tpr_m4, color=roc_colors[3], linewidth=2,

```



```


        label=f'Decision Tree (AUC = {macro_auc_m4:.2f})')
auc_scores_list.append({'Model': 'Decision Tree', 'AUC': macro_auc_m4})

# --- Bagging DT ---
y_score_m5 = m5.predict_proba(X_test_scaled)
fpr_m5, tpr_m5, roc_auc_m5 = {}, {}, {}
for i in range(n_classes):
    fpr_m5[i], tpr_m5[i], _ = roc_curve(y_test_bin[:, i], y_score_m5[:, i])
    roc_auc_m5[i] = auc(fpr_m5[i], tpr_m5[i])
all_fpr_m5 = np.unique(np.concatenate([fpr_m5[i] for i in range(n_classes)]))
mean_tpr_m5 = np.zeros_like(all_fpr_m5)
for i in range(n_classes):
    mean_tpr_m5 += np.interp(all_fpr_m5, fpr_m5[i], tpr_m5[i])
mean_tpr_m5 /= n_classes
macro_auc_m5 = auc(all_fpr_m5, mean_tpr_m5)
plt.plot(all_fpr_m5, mean_tpr_m5, color=roc_colors[4], linewidth=2,
        label=f'Bagging DT (AUC = {macro_auc_m5:.2f})')
auc_scores_list.append({'Model': 'Bagging DT', 'AUC': macro_auc_m5})

# --- XGBoost ---
y_score_xgb = m_xgb.predict_proba(X_test_scaled) # Predictions for XGBoost
fpr_xgb, tpr_xgb, roc_auc_xgb = {}, {}, {}
for i in range(n_classes):
    fpr_xgb[i], tpr_xgb[i], _ = roc_curve(y_test_bin[:, i], y_score_xgb[:, i])
    roc_auc_xgb[i] = auc(fpr_xgb[i], tpr_xgb[i])
all_fpr_xgb = np.unique(np.concatenate([fpr_xgb[i] for i in range(n_classes)]))
mean_tpr_xgb = np.zeros_like(all_fpr_xgb)
for i in range(n_classes):
    mean_tpr_xgb += np.interp(all_fpr_xgb, fpr_xgb[i], tpr_xgb[i])
mean_tpr_xgb /= n_classes
macro_auc_xgb = auc(all_fpr_xgb, mean_tpr_xgb)
plt.plot(all_fpr_xgb, mean_tpr_xgb, color=roc_colors[5], linewidth=2,
        label=f'XGBoost (AUC = {macro_auc_xgb:.2f})')
auc_scores_list.append({'Model': 'XGBoost', 'AUC': macro_auc_xgb})

plt.plot([0, 1], [0, 1], 'k--', linewidth=1, label='Random Guessing')
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('Macro-Averaged ROC Curves - All Models',
        fontsize=13, fontweight='bold')
plt.legend(loc='lower right', fontsize=9)
plt.grid(True, linestyle='--', alpha=0.5)
plt.tight_layout()
plt.savefig('plot8_roc_curves.png', dpi=150)
plt.show()
print('Saved: plot8_roc_curves.png')

```



Saved: plot8_roc_curves.png

```
In [289]: auc_vals_dict = {
    'Logistic Regression': macro_auc_m1,
    'Logit L1'           : macro_auc_m2,
    'Logit L2'           : macro_auc_m3,
    'Decision Tree'      : macro_auc_m4,
    'Bagging DT'         : macro_auc_m5,
    'XGBoost'            : macro_auc_xgb # Added XGBoost AUC
}

summary_df = pd.DataFrame({
    'Model'       : model_names,
    'Precision'    : precision_macro_scores,
    'Recall'       : recall_macro_scores,
    'F1-Score'     : f1_macro_scores,
    'AUC'          : [macro_auc_m1, macro_auc_m2, macro_auc_m3, macro_auc_m4, macro_auc_m5, macro_auc_xgb],
    'Accuracy'     : accuracy_scores
})

summary_df = summary_df.sort_values('F1-Score', ascending=False).reset_index(drop=True)
summary_df.index += 1

print('=' * 65)
print('  FINAL RESULTS - Ranked by Macro F1-Score')
print('=' * 65)
display(summary_df)
```

=====

FINAL RESULTS – Ranked by Macro F1-Score

=====

	Model	Precision	Recall	F1-Score	AUC	Accuracy
1	XGBoost	0.961230	0.956258	0.958598	0.998877	0.955
2	Logistic Regression	0.961400	0.950071	0.954576	0.996385	0.955
3	Logit L1	0.945012	0.937838	0.940833	0.995876	0.940
4	Logit L2	0.925307	0.897222	0.907773	0.980754	0.910
5	Bagging DT	0.914488	0.899364	0.905934	0.987020	0.900
6	Decision Tree	0.906018	0.891605	0.897761	0.961525	0.890